

1 CKAD

- 1.1 Table of Contents
- 1.2 Overview
- 1.3 Exam Overview
- 1.4 Exam Domains & Weights
- 1.5 Prerequisites
- 1.6 Study Resources
 - 1.6.1 Official Resources
 - 1.6.2 Recommended Courses
 - 1.6.3 Practice Resources
- 1.7 Quick Navigation
- 1.8 Exam Environment
 - 1.8.1 Allowed Resources During Exam
- 1.9 Exam Tips
- 1.10 Useful kubectl Commands
- 1.11 Registration
- 1.12 Application Design and Build
- 1.13 Container Images
 - 1.13.1 Building Container Images
 - 1.13.2 Multi-stage Builds
 - 1.13.3 Image Best Practices
- 1.14 Jobs and CronJobs
 - 1.14.1 Job
 - 1.14.2 CronJob
- 1.15 Multi-Container Pods
 - 1.15.1 Sidecar Pattern
 - 1.15.2 Init Containers
 - 1.15.3 Ambassador Pattern
 - 1.15.4 Adapter Pattern
- 1.16 Volumes
 - 1.16.1 emptyDir
 - 1.16.2 hostPath
 - 1.16.3 PersistentVolumeClaim
- 1.17 Kubectl Imperative Commands
 - 1.17.1 Create Resources Quickly
- 1.18 Key Concepts to Remember
- 1.19 Practice Questions
- 1.20 Application Deployment
- 1.21 Deployments
 - 1.21.1 Creating Deployments
 - 1.21.2 Imperative Commands
- 1.22 Deployment Strategies
 - 1.22.1 Rolling Update (Default)
 - 1.22.2 Recreate
- 1.23 Rollouts
 - 1.23.1 Rollout Commands
 - 1.23.2 Recording Changes
- 1.24 Blue-Green Deployment
- 1.25 Canary Deployment
- 1.26 Helm
 - 1.26.1 Helm Basics

- 1.26.2 Helm Chart Structure
- 1.26.3 Chart.yaml
- 1.26.4 values.yaml
- 1.26.5 Template Example
- 1.27 Kustomize
 - 1.27.1 Kustomize Structure
 - 1.27.2 Base kustomization.yaml
 - 1.27.3 Overlay kustomization.yaml
 - 1.27.4 Kustomize Commands
- 1.28 Key Concepts to Remember
- 1.29 Practice Questions
- 1.30 Application Observability and Maintenance
- 1.31 Probes
 - 1.31.1 Liveness Probe
 - 1.31.2 Readiness Probe
 - 1.31.3 Startup Probe
 - 1.31.4 Probe Types
 - 1.31.5 Probe Parameters
- 1.32 Logging
 - 1.32.1 Viewing Logs
 - 1.32.2 Logging Architecture
- 1.33 Debugging
 - 1.33.1 Debug Commands
 - 1.33.2 Common Issues
 - 1.33.3 Pod Status Phases
- 1.34 Monitoring
 - 1.34.1 Resource Metrics
 - 1.34.2 Metrics Server
- 1.35 Application Maintenance
 - 1.35.1 Updating Applications
 - 1.35.2 Scaling
 - 1.35.3 HorizontalPodAutoscaler
- 1.36 Key Concepts to Remember
- 1.37 Practice Questions
- 1.38 Application Environment, Configuration and Security
- 1.39 ConfigMaps
 - 1.39.1 Creating ConfigMaps
 - 1.39.2 ConfigMap YAML
 - 1.39.3 Using ConfigMaps
- 1.40 Secrets
 - 1.40.1 Creating Secrets
 - 1.40.2 Secret YAML
 - 1.40.3 Secret Types
 - 1.40.4 Using Secrets
- 1.41 ServiceAccounts
 - 1.41.1 Creating ServiceAccounts
 - 1.41.2 ServiceAccount YAML
 - 1.41.3 Using ServiceAccounts
- 1.42 Security Context
 - 1.42.1 Pod-level Security Context
 - 1.42.2 Container-level Security Context
 - 1.42.3 Security Context Fields
- 1.43 Resource Requirements
 - 1.43.1 Requests and Limits

- 1.43.2 CPU Units
- 1.43.3 Memory Units
- 1.43.4 QoS Classes
- 1.44 LimitRange
- 1.45 ResourceQuota
- 1.46 Admission Controllers
 - 1.46.1 Pod Security Standards
- 1.47 Key Concepts to Remember
- 1.48 Practice Questions
- 1.49 Services and Networking
- 1.50 Services
 - 1.50.1 Service Types
 - 1.50.2 ClusterIP Service
 - 1.50.3 NodePort Service
 - 1.50.4 LoadBalancer Service
 - 1.50.5 ExternalName Service
 - 1.50.6 Headless Service
 - 1.50.7 Creating Services Imperatively
- 1.51 DNS in Kubernetes
 - 1.51.1 Service DNS
 - 1.51.2 Pod DNS
 - 1.51.3 DNS Resolution Example
- 1.52 Network Policies
 - 1.52.1 Default Behavior
 - 1.52.2 Deny All Ingress
 - 1.52.3 Deny All Egress
 - 1.52.4 Allow Specific Ingress
 - 1.52.5 Allow Egress to Specific Pods
 - 1.52.6 Network Policy Selectors
- 1.53 Ingress
 - 1.53.1 Ingress Resource
 - 1.53.2 Multiple Hosts
 - 1.53.3 Path-based Routing
 - 1.53.4 TLS Ingress
 - 1.53.5 Path Types
 - 1.53.6 Creating Ingress Imperatively
- 1.54 Port Forwarding
- 1.55 Key Concepts to Remember
- 1.56 Practice Questions
- 1.57 Sample Practice Questions
- 1.58 Practice Resources
- 1.59 Instructions
- 1.60 Section 1: Application Design and Build (20%)
 - 1.60.1 Question 1.1 - Create a Job
 - 1.60.2 Question 1.2 - Create a CronJob
 - 1.60.3 Question 1.3 - Multi-Container Pod
- 1.61 Section 2: Application Deployment (20%)
 - 1.61.1 Question 2.1 - Create and Scale Deployment
 - 1.61.2 Question 2.2 - Deployment Strategy
 - 1.61.3 Question 2.3 - Helm Operations
- 1.62 Section 3: Application Observability and Maintenance (15%)
 - 1.62.1 Question 3.1 - Configure Probes
 - 1.62.2 Question 3.2 - Debugging
 - 1.62.3 Question 3.3 - View Logs

- 1.63 Section 4: Application Environment, Configuration and Security (25%)
 - 1.63.1 Question 4.1 - ConfigMap and Secret
 - 1.63.2 Question 4.2 - Security Context
 - 1.63.3 Question 4.3 - Resource Limits
- 1.64 Section 5: Services and Networking (20%)
 - 1.64.1 Question 5.1 - Create Services
 - 1.64.2 Question 5.2 - Network Policy
 - 1.64.3 Question 5.3 - Ingress
- 1.65 Exam Tips
- 1.66 Additional Practice

1 CKAD

Generated on: 2026-01-13 15:04:29 Version: 1.0

1.1 Table of Contents

1. [Overview](#)
 2. [Application Design and Build](#)
 3. [Application Deployment](#)
 4. [Application Observability and Maintenance](#)
 5. [Application Environment, Configuration and Security](#)
 6. [Services and Networking](#)
 7. [Sample Practice Questions](#)
-

1.2 Overview



The **Certified Kubernetes Application Developer (CKAD)** exam certifies that candidates can design, build, and deploy cloud native applications for Kubernetes.

1.3 Exam Overview

| Detail | Information |
|-------------------------------|------------------------------|
| Exam Format | Performance-based (hands-on) |
| Number of Questions | 15-20 |
| Duration | 2 hours |
| Passing Score | 66% |
| Certification Validity | 3 years |
| Cost | \$395 USD |

| Detail | Information |
|--------------------|---------------|
| Retake Policy | 1 free retake |
| Kubernetes Version | 1.30 |

1.4 Exam Domains & Weights

| Domain | Weight |
|---|--------|
| Application Design and Build | 20% |
| Application Deployment | 20% |
| Application Observability and Maintenance | 15% |
| Application Environment, Configuration and Security | 25% |
| Services and Networking | 20% |

1.5 Prerequisites

- Basic understanding of Kubernetes concepts
- Familiarity with YAML syntax
- Linux command line proficiency
- Container fundamentals (Docker)

1.6 Study Resources

1.6.1 Official Resources

- [CKAD Exam Curriculum](#)
- [Kubernetes Documentation](#)
- [Kubernetes Tasks](#)

1.6.2 Recommended Courses

- [Kubernetes for Developers \(LFD259\)](#)
- [CKAD with Tests - Udemy](#)

1.6.3 Practice Resources

- [Killercoda CKAD Scenarios](#) ★ Highly Recommended
- [Kubernetes Playground](#)
- [killer.sh CKAD Simulator](#)

1.7 Quick Navigation

- [01 - Application Design and Build](#)
- [02 - Application Deployment](#)
- [03 - Application Observability and Maintenance](#)
- [04 - Application Environment, Configuration and Security](#)
- [05 - Services and Networking](#)
- [Sample Practice Questions](#)

1.8 Exam Environment

The CKAD exam provides:

- Access to multiple Kubernetes clusters
- `kubectl` with auto-completion enabled
- Access to Kubernetes documentation (kubernetes.io)
- A Linux terminal environment

1.8.1 Allowed Resources During Exam

- kubernetes.io/docs
- kubernetes.io/blog
- helm.sh/docs

1.9 Exam Tips

1. **Practice `kubectl` imperative commands** - They save significant time
2. **Master YAML generation** - Use `kubectl run --dry-run=client -o yaml`
3. **Know vim/nano basics** - You'll edit YAML files frequently
4. **Use aliases** - Set up alias `k=kubectl` and enable auto-completion
5. **Bookmark important docs** - Prepare bookmarks for quick access
6. **Time management** - Don't spend too long on any single question
7. **Practice on Killercoda** - Free hands-on scenarios at killercoda.com/ckad

1.10 Useful `kubectl` Commands

```
# Set alias
alias k=kubectl

# Enable auto-completion
source <(kubectl completion bash)
complete -o default -F __start_kubectl k

# Generate YAML templates
k run nginx --image=nginx --dry-run=client -o yaml > pod.yaml
k create deployment nginx --image=nginx --dry-run=client -o yaml
> deploy.yaml
```

```
k create service clusterip nginx --tcp=80:80 --dry-run=client -o
yaml > svc.yaml

# Quick pod creation
k run nginx --image=nginx --port=80 --labels=app=web

# Expose deployment
k expose deployment nginx --port=80 --target-port=80 --
type=ClusterIP

# Scale deployment
k scale deployment nginx --replicas=3

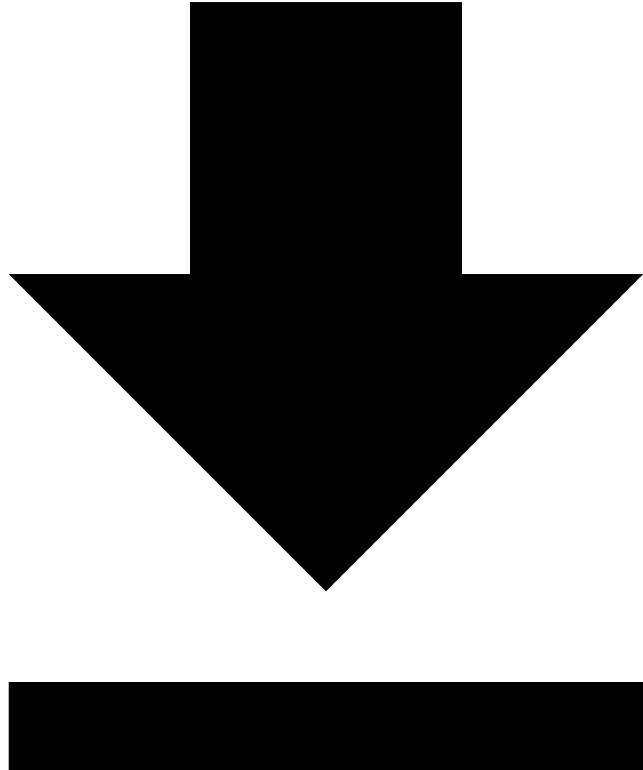
# Set resources
k set resources deployment nginx --limits=cpu=200m, memory=512Mi

# Rollout commands
k rollout status deployment/nginx
k rollout history deployment/nginx
k rollout undo deployment/nginx
```

1.11 Registration

[Register for CKAD Exam](#)

1.12 Application Design and Build



[Download PDF Version](#)

This domain covers designing and building cloud native applications for Kubernetes.

1.13 Container Images

1.13.1 Building Container Images

```
# Example Dockerfile
FROM python:3.11-slim

WORKDIR /app

COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt

COPY . .
```

```
EXPOSE 8080  
  
USER 1000  
  
CMD ["python", "app.py"]
```

1.13.2 Multi-stage Builds

```
# Build stage  
FROM golang:1.21 AS builder  
WORKDIR /app  
COPY . .  
RUN CGO_ENABLED=0 go build -o myapp  
  
# Runtime stage  
FROM alpine:3.18  
COPY --from=builder /app/myapp /myapp  
ENTRYPOINT ["/myapp"]
```

1.13.3 Image Best Practices

- Use specific image tags, not latest
- Use minimal base images (alpine, distroless)
- Run as non-root user
- Use multi-stage builds to reduce image size
- Scan images for vulnerabilities

1.14 Jobs and CronJobs

1.14.1 Job

A Job creates one or more Pods and ensures they complete successfully.

```
apiVersion: batch/v1  
kind: Job  
metadata:  
  name: pi-job  
spec:  
  completions: 3  
  parallelism: 2  
  backoffLimit: 4  
  activeDeadlineSeconds: 100  
  template:  
    spec:  
      containers:  
      - name: pi  
        image: perl:5.34
```

```

  command: ["perl", "-Mbignum=bpi", "-wle", "print
    bpi(2000)"]
  restartPolicy: Never

```

Key Fields:

| Field | Description |
|-----------------------|--|
| completions | Number of successful completions required |
| parallelism | Number of pods running in parallel |
| backoffLimit | Number of retries before marking as failed |
| activeDeadlineSeconds | Maximum time for the job |
| restartPolicy | Must be Never or OnFailure |

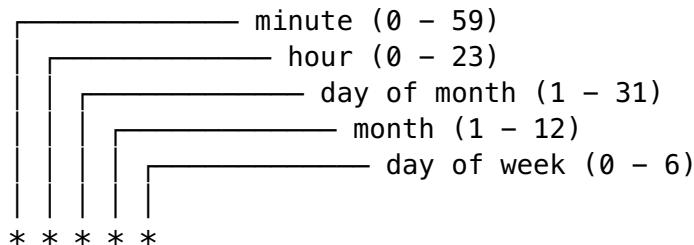
1.14.2 CronJob

```

apiVersion: batch/v1
kind: CronJob
metadata:
  name: backup-cronjob
spec:
  schedule: "0 2 * * *" # Daily at 2 AM
  concurrencyPolicy: Forbid
  successfulJobsHistoryLimit: 3
  failedJobsHistoryLimit: 1
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: backup
              image: backup-tool:v1
              command: ["/bin/sh", "-c", "backup.sh"]
              restartPolicy: OnFailure

```

Cron Schedule Format:



Concurrency Policies:

| Policy | Description |
|--------|---|
| Allow | Allow concurrent jobs (default) |
| Forbid | Skip new job if previous is still running |

| Policy | Description |
|---------|--|
| Replace | Replace currently running job with new one |

1.15 Multi-Container Pods

1.15.1 Sidecar Pattern

```
apiVersion: v1
kind: Pod
metadata:
  name: app-with-sidecar
spec:
  containers:
    - name: app
      image: myapp:v1
      volumeMounts:
        - name: logs
          mountPath: /var/log/app
    - name: log-shipper
      image: fluentd:v1
      volumeMounts:
        - name: logs
          mountPath: /var/log/app
  volumes:
    - name: logs
      emptyDir: {}
```

1.15.2 Init Containers

Init containers run before app containers start:

```
apiVersion: v1
kind: Pod
metadata:
  name: app-with-init
spec:
  initContainers:
    - name: init-db
      image: busybox:1.36
      command: ['sh', '-c', 'until nc -z db-service 5432; do sleep 2; done']
    - name: init-config
      image: busybox:1.36
      command: ['sh', '-c', 'wget -O /config/app.conf http://config-server/app.conf']
      volumeMounts:
        - name: config
          mountPath: /config
  containers:
    - name: app
```

```
image: myapp:v1
volumeMounts:
- name: config
  mountPath: /config
volumes:
- name: config
  emptyDir: {}
```

1.15.3 Ambassador Pattern

```
apiVersion: v1
kind: Pod
metadata:
  name: app-with-ambassador
spec:
  containers:
  - name: app
    image: myapp:v1
    env:
    - name: DB_HOST
      value: "localhost"
    - name: DB_PORT
      value: "5432"
  - name: ambassador
    image: ambassador-proxy:v1
    ports:
    - containerPort: 5432
```

1.15.4 Adapter Pattern

```
apiVersion: v1
kind: Pod
metadata:
  name: app-with-adapter
spec:
  containers:
  - name: app
    image: legacy-app:v1
    volumeMounts:
    - name: logs
      mountPath: /var/log
  - name: adapter
    image: log-adapter:v1
    volumeMounts:
    - name: logs
      mountPath: /var/log
```

1.16 Volumes

1.16.1 emptyDir

Temporary storage that exists for the Pod's lifetime:

```
volumes:  
- name: cache  
  emptyDir: {}  
- name: memory-cache  
  emptyDir:  
    medium: Memory  
    sizeLimit: 100Mi
```

1.16.2 hostPath

Mount a file or directory from the host node:

```
volumes:  
- name: host-data  
  hostPath:  
    path: /data  
    type: DirectoryOrCreate
```

1.16.3 PersistentVolumeClaim

```
volumes:  
- name: data  
  persistentVolumeClaim:  
    claimName: my-pvc
```

1.17 Kubectl Imperative Commands

1.17.1 Create Resources Quickly

```
# Create a pod  
kubectl run nginx --image=nginx --port=80  
  
# Create a deployment  
kubectl create deployment nginx --image=nginx --replicas=3  
  
# Create a job  
kubectl create job pi --image=perl -- perl -Mbignum=bpi -wle  
  'print bpi(2000)'  
  
# Create a cronjob  
kubectl create cronjob backup --image=backup --schedule="0 2 * * *"  
  -- /bin/sh -c 'backup.sh'
```

```
# Generate YAML
kubectl run nginx --image=nginx --dry-run=client -o yaml >
pod.yaml
```

1.18 Key Concepts to Remember

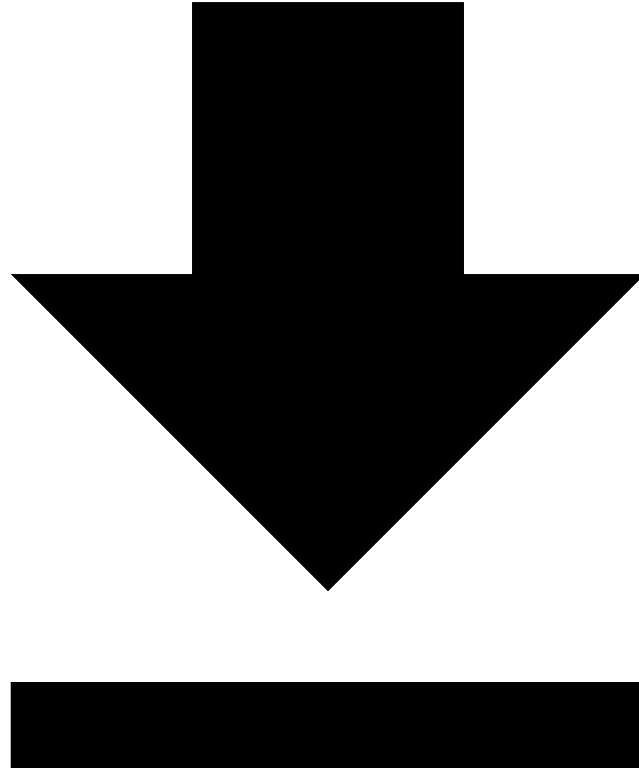
1. **Jobs** - Run to completion, use `restartPolicy: Never` or `OnFailure`
2. **CronJobs** - Scheduled jobs using cron syntax
3. **Init containers** - Run before main containers, must complete successfully
4. **Sidecar pattern** - Helper container alongside main app
5. **Multi-stage builds** - Reduce image size

1.19 Practice Questions

1. How do you create a Job that runs 5 completions with 2 parallel pods?
2. What is the difference between init containers and sidecar containers?
3. How do you generate YAML for a pod without creating it?
4. What cron schedule runs every Monday at 3 AM?
5. What happens if a CronJob's concurrencyPolicy is set to Forbid?

[Back to CKAD Overview](#) | [Next: Application Deployment →](#)

1.20 Application Deployment



[Download PDF Version](#)

This domain covers deploying applications in Kubernetes using various strategies and tools.

1.21 Deployments

1.21.1 Creating Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
```

```

    app: nginx
  template:
    metadata:
      labels:
        app: nginx
  spec:
    containers:
      - name: nginx
        image: nginx:1.21
        ports:
          - containerPort: 80
        resources:
          requests:
            memory: "64Mi"
            cpu: "250m"
          limits:
            memory: "128Mi"
            cpu: "500m"

```

1.21.2 Imperative Commands

```

# Create deployment
kubectl create deployment nginx --image=nginx:1.21 --replicas=3

# Scale deployment
kubectl scale deployment nginx --replicas=5

# Update image
kubectl set image deployment/nginx nginx=nginx:1.22

# Edit deployment
kubectl edit deployment nginx

```

1.22 Deployment Strategies

1.22.1 Rolling Update (Default)

```

spec:
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 25%
      maxUnavailable: 25%

```

| Parameter | Description |
|----------------|--|
| maxSurge | Max pods above desired count during update |
| maxUnavailable | Max pods unavailable during update |

1.22.2 Recreate

```
spec:  
  strategy:  
    type: Recreate
```

All existing pods are killed before new ones are created.

1.23 Rollouts

1.23.1 Rollout Commands

```
# Check rollout status  
kubectl rollout status deployment/nginx  
  
# View rollout history  
kubectl rollout history deployment/nginx  
  
# View specific revision  
kubectl rollout history deployment/nginx --revision=2  
  
# Undo rollout (rollback)  
kubectl rollout undo deployment/nginx  
  
# Rollback to specific revision  
kubectl rollout undo deployment/nginx --to-revision=2  
  
# Pause rollout  
kubectl rollout pause deployment/nginx  
  
# Resume rollout  
kubectl rollout resume deployment/nginx  
  
# Restart deployment  
kubectl rollout restart deployment/nginx
```

1.23.2 Recording Changes

```
# Record the command in revision history  
kubectl set image deployment/nginx nginx=nginx:1.22 --record
```

1.24 Blue-Green Deployment

Manual blue-green deployment using Services:

```
# Blue deployment (current)  
apiVersion: apps/v1  
kind: Deployment  
metadata:
```

```

    name: app-blue
spec:
  replicas: 3
  selector:
    matchLabels:
      app: myapp
      version: blue
  template:
    metadata:
      labels:
        app: myapp
        version: blue
    spec:
      containers:
        - name: app
          image: myapp:v1
  ---
# Green deployment (new)
apiVersion: apps/v1
kind: Deployment
metadata:
  name: app-green
spec:
  replicas: 3
  selector:
    matchLabels:
      app: myapp
      version: green
  template:
    metadata:
      labels:
        app: myapp
        version: green
    spec:
      containers:
        - name: app
          image: myapp:v2
  ---
# Service - switch selector to change versions
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  selector:
    app: myapp
    version: blue # Change to 'green' to switch
  ports:
    - port: 80
      targetPort: 8080

```

1.25 Canary Deployment

```
# Stable deployment (90% traffic)
apiVersion: apps/v1
kind: Deployment
metadata:
  name: app-stable
spec:
  replicas: 9
  selector:
    matchLabels:
      app: myapp
  template:
    metadata:
      labels:
        app: myapp
    spec:
      containers:
        - name: app
          image: myapp:v1
---
# Canary deployment (10% traffic)
apiVersion: apps/v1
kind: Deployment
metadata:
  name: app-canary
spec:
  replicas: 1
  selector:
    matchLabels:
      app: myapp
  template:
    metadata:
      labels:
        app: myapp
    spec:
      containers:
        - name: app
          image: myapp:v2
---
# Service routes to both (based on replica ratio)
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  selector:
    app: myapp # Matches both deployments
  ports:
    - port: 80
      targetPort: 8080
```

1.26 Helm

1.26.1 Helm Basics

```
# Add repository
helm repo add bitnami https://charts.bitnami.com/bitnami
helm repo update

# Search charts
helm search repo nginx
helm search hub wordpress

# Install chart
helm install my-nginx bitnami/nginx

# Install with custom values
helm install my-nginx bitnami/nginx -f values.yaml
helm install my-nginx bitnami/nginx --set service.type=NodePort

# List releases
helm list
helm list -A # All namespaces

# Upgrade release
helm upgrade my-nginx bitnami/nginx --set replicaCount=3

# Rollback
helm rollback my-nginx 1

# Uninstall
helm uninstall my-nginx

# Show chart info
helm show values bitnami/nginx
helm show chart bitnami/nginx
```

1.26.2 Helm Chart Structure

```
mychart/
├── Chart.yaml      # Chart metadata
├── values.yaml    # Default configuration values
├── charts/         # Chart dependencies
└── templates/      # Template files
    ├── deployment.yaml
    ├── service.yaml
    ├── _helpers.tpl   # Template helpers
    └── NOTES.txt     # Post-install notes
└── .helmignore     # Files to ignore
```

1.26.3 Chart.yaml

```
apiVersion: v2
name: mychart
description: A Helm chart for my application
type: application
version: 0.1.0
appVersion: "1.0.0"
```

1.26.4 values.yaml

```
replicaCount: 3

image:
  repository: nginx
  tag: "1.21"
  pullPolicy: IfNotPresent

service:
  type: ClusterIP
  port: 80

resources:
  limits:
    cpu: 100m
    memory: 128Mi
  requests:
    cpu: 100m
    memory: 128Mi
```

1.26.5 Template Example

```
# templates/deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: {{ .Release.Name }}-deployment
spec:
  replicas: {{ .Values.replicaCount }}
  selector:
    matchLabels:
      app: {{ .Release.Name }}
  template:
    metadata:
      labels:
        app: {{ .Release.Name }}
  spec:
    containers:
    - name: {{ .Chart.Name }}
      image: "{{ .Values.image.repository }}:
{{ .Values.image.tag }}"
```

```
  ports:  
    - containerPort: 80
```

1.27 Kustomize

1.27.1 Kustomize Structure

```
└── base/  
    ├── deployment.yaml  
    ├── service.yaml  
    └── kustomization.yaml  
└── overlays/  
    ├── dev/  
    │   └── kustomization.yaml  
    └── prod/  
        └── kustomization.yaml
```

1.27.2 Base kustomization.yaml

```
# base/kustomization.yaml  
apiVersion: kustomize.config.k8s.io/v1beta1  
kind: Kustomization  
  
resources:  
  - deployment.yaml  
  - service.yaml
```

1.27.3 Overlay kustomization.yaml

```
# overlays/prod/kustomization.yaml  
apiVersion: kustomize.config.k8s.io/v1beta1  
kind: Kustomization  
  
resources:  
  - ../../base  
  
namePrefix: prod-  
  
replicas:  
  - name: myapp  
    count: 5  
  
images:  
  - name: myapp  
    newTag: v2.0.0  
  
patches:  
  - patch: |-  
    - op: replace
```

```
path: /spec/template/spec/containers/0/resources/limits/
  memory
    value: 512Mi
  target:
    kind: Deployment
    name: myapp
```

1.27.4 Kustomize Commands

```
# Preview output
kubectl kustomize overlays/prod/

# Apply
kubectl apply -k overlays/prod/

# Delete
kubectl delete -k overlays/prod/
```

1.28 Key Concepts to Remember

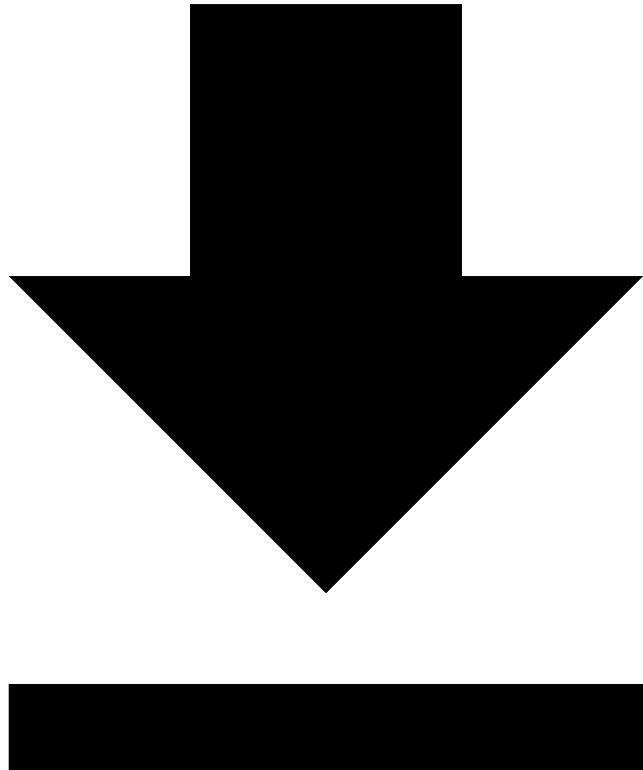
1. **Rolling Update** - Default strategy, gradual replacement
2. **Recreate** - Kill all, then create new
3. **Rollout commands** - status, history, undo, pause, resume
4. **Helm** - Package manager for Kubernetes
5. **Kustomize** - Template-free configuration customization

1.29 Practice Questions

1. How do you rollback a deployment to revision 3?
2. What is the difference between maxSurge and maxUnavailable?
3. How do you install a Helm chart with custom values?
4. What command shows the rollout history of a deployment?
5. How do you apply a Kustomize overlay?

[← Previous: Application Design and Build](#) | [Back to CKAD Overview](#) | [Next: Application Observability and Maintenance →](#)

1.30 Application Observability and Maintenance



[Download PDF Version](#)

This domain covers monitoring, debugging, and maintaining applications in Kubernetes.

1.31 Probes

1.31.1 Liveness Probe

Determines if a container is running. If it fails, the container is restarted.

```
apiVersion: v1
kind: Pod
metadata:
  name: app-with-liveness
spec:
  containers:
  - name: app
    image: myapp:v1
```

```
livenessProbe:  
  httpGet:  
    path: /healthz  
    port: 8080  
  initialDelaySeconds: 15  
  periodSeconds: 10  
  timeoutSeconds: 5  
  failureThreshold: 3  
  successThreshold: 1
```

1.31.2 Readiness Probe

Determines if a container is ready to receive traffic.

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: app-with-readiness  
spec:  
  containers:  
  - name: app  
    image: myapp:v1  
    readinessProbe:  
      httpGet:  
        path: /ready  
        port: 8080  
      initialDelaySeconds: 5  
      periodSeconds: 5  
      failureThreshold: 3
```

1.31.3 Startup Probe

Used for slow-starting containers. Disables liveness/readiness until it succeeds.

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: app-with-startup  
spec:  
  containers:  
  - name: app  
    image: myapp:v1  
    startupProbe:  
      httpGet:  
        path: /healthz  
        port: 8080  
      failureThreshold: 30  
      periodSeconds: 10  
    livenessProbe:  
      httpGet:  
        path: /healthz
```

```

    port: 8080
    periodSeconds: 10

```

1.31.4 Probe Types

| Type | Description |
|-----------|---|
| httpGet | HTTP GET request to specified path and port |
| tcpSocket | TCP connection to specified port |
| exec | Execute command in container |
| grpc | gRPC health check |

```

# TCP Socket probe
livenessProbe:
  tcpSocket:
    port: 3306
    initialDelaySeconds: 15
    periodSeconds: 10

# Exec probe
livenessProbe:
  exec:
    command:
      - cat
      - /tmp/healthy
    initialDelaySeconds: 5
    periodSeconds: 5

# gRPC probe
livenessProbe:
  grpc:
    port: 50051
    initialDelaySeconds: 10

```

1.31.5 Probe Parameters

| Parameter | Description | Default |
|---------------------|------------------------------------|---------|
| initialDelaySeconds | Delay before first probe | 0 |
| periodSeconds | How often to probe | 10 |
| timeoutSeconds | Probe timeout | 1 |
| failureThreshold | Failures before action | 3 |
| successThreshold | Successes to be considered healthy | 1 |

1.32 Logging

1.32.1 Viewing Logs

```
# View pod logs
kubectl logs nginx

# View specific container logs
kubectl logs nginx -c sidecar

# Follow logs
kubectl logs -f nginx

# View previous container logs (after restart)
kubectl logs nginx --previous

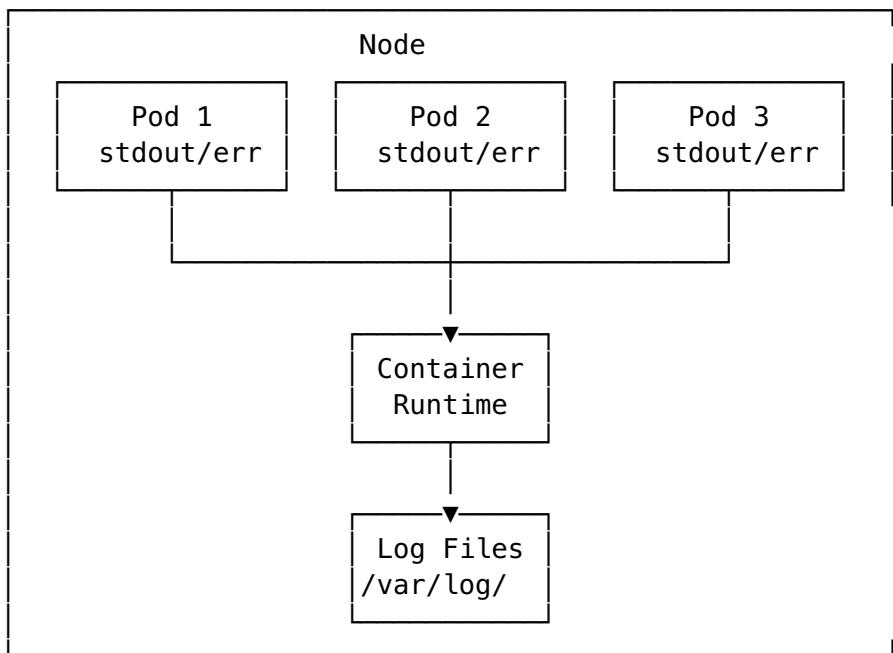
# View last N lines
kubectl logs nginx --tail=100

# View logs since time
kubectl logs nginx --since=1h
kubectl logs nginx --since-time=2024-01-01T00:00:00Z

# View logs from all pods with label
kubectl logs -l app=nginx

# View logs from all containers in pod
kubectl logs nginx --all-containers
```

1.32.2 Logging Architecture



1.33 Debugging

1.33.1 Debug Commands

```
# Describe pod (events, status)
kubectl describe pod nginx

# Get pod details
kubectl get pod nginx -o yaml
kubectl get pod nginx -o wide

# Check events
kubectl get events --sort-by=.lastTimestamp'
kubectl get events --field-selector involvedObject.name=nginx

# Execute command in container
kubectl exec nginx -- ls /app
kubectl exec -it nginx -- /bin/sh

# Copy files to/from container
kubectl cp nginx:/var/log/app.log ./app.log
kubectl cp ./config.yaml nginx:/app/config.yaml

# Port forward
kubectl port-forward pod/nginx 8080:80
kubectl port-forward svc/nginx 8080:80

# Debug with ephemeral container
kubectl debug nginx -it --image=busybox --target=nginx
```

1.33.2 Common Issues

| Issue | Debug Steps |
|-----------------------------------|--|
| ImagePullBackOff | Check image name, registry access, pull secrets |
| CrashLoopBackOff | Check logs, probe configuration, resource limits |
| Pending | Check events, node resources, taints/tolerations |
| OOMKilled | Increase memory limits |
| CreateContainerConfigError | Check ConfigMaps, Secrets references |

1.33.3 Pod Status Phases

| Phase | Description |
|-----------|--|
| Pending | Pod accepted but not running |
| Running | Pod bound to node, containers running |
| Succeeded | All containers terminated successfully |
| Failed | All containers terminated, at least one failed |

| Phase | Description |
|---------|--------------------------------|
| Unknown | Pod state cannot be determined |

1.34 Monitoring

1.34.1 Resource Metrics

```
# View node resource usage
kubectl top nodes

# View pod resource usage
kubectl top pods
kubectl top pods -A
kubectl top pods --containers

# Sort by CPU/memory
kubectl top pods --sort-by=cpu
kubectl top pods --sort-by=memory
```

1.34.2 Metrics Server

Required for kubectl top commands:

```
# Check if metrics server is running
kubectl get pods -n kube-system | grep metrics-server

# Install metrics server (if needed)
kubectl apply -f https://github.com/kubernetes-sigs/metrics-
server/releases/latest/download/components.yaml
```

1.35 Application Maintenance

1.35.1 Updating Applications

```
# Update image
kubectl set image deployment/nginx nginx=nginx:1.22

# Update environment variable
kubectl set env deployment/nginx ENV=production

# Update resources
kubectl set resources deployment/nginx --
  limits=cpu=200m,memory=512Mi

# Patch resource
kubectl patch deployment nginx -p '{"spec":{"replicas":5}}'
```

1.35.2 Scaling

```
# Manual scaling
kubectl scale deployment nginx --replicas=5

# Autoscaling
kubectl autoscale deployment nginx --min=2 --max=10 --cpu-
    percent=80
```

1.35.3 HorizontalPodAutoscaler

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: nginx-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: nginx
  minReplicas: 2
  maxReplicas: 10
  metrics:
  - type: Resource
    resource:
      name: cpu
      target:
        type: Utilization
        averageUtilization: 80
  - type: Resource
    resource:
      name: memory
      target:
        type: Utilization
        averageUtilization: 80
```

1.36 Key Concepts to Remember

1. **Liveness** - Is the container running? Restart if not
2. **Readiness** - Is the container ready for traffic?
3. **Startup** - For slow-starting containers
4. **kubectl logs** - View container output
5. **kubectl describe** - Detailed resource info with events
6. **kubectl top** - Resource usage (requires metrics-server)

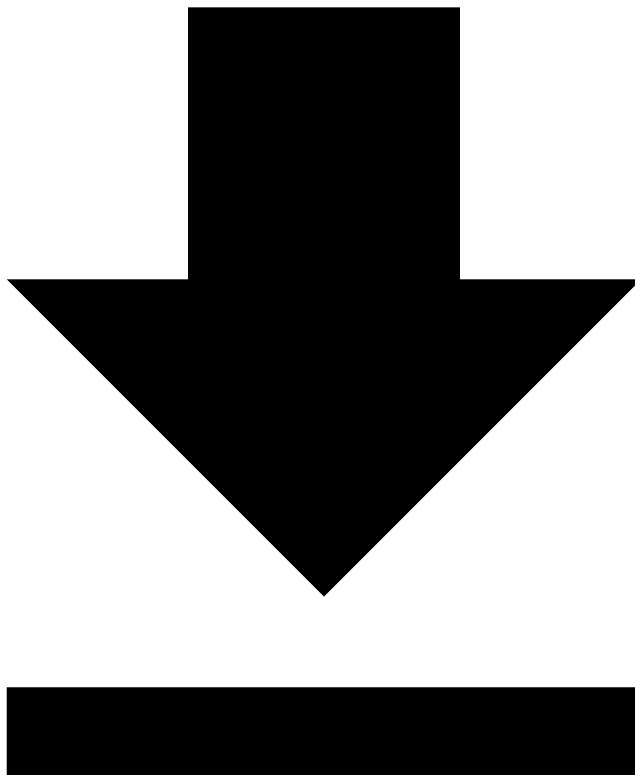
1.37 Practice Questions

1. What happens when a liveness probe fails?
2. How do you view logs from a previous container instance?
3. What is the difference between readiness and liveness probes?

4. How do you execute a command in a running container?
 5. What probe type would you use for a database container?
-

[← Previous: Application Deployment](#) | [Back to CKAD Overview](#) | [Next: Application Environment, Configuration and Security →](#)

1.38 Application Environment, Configuration and Security



[Download PDF Version](#)

This domain covers configuring applications, managing secrets, and implementing security in Kubernetes.

1.39 ConfigMaps

1.39.1 Creating ConfigMaps

```
# From literal values
kubectl create configmap app-config --from-literal=ENV=production
--from-literal=LOG_LEVEL=info

# From file
kubectl create configmap app-config --from-file=config.properties

# From directory
kubectl create configmap app-config --from-file=config/

# From env file
kubectl create configmap app-config --from-env-file=app.env
```

1.39.2 ConfigMap YAML

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-config
data:
  ENV: production
  LOG_LEVEL: info
  config.json: |
    {
      "database": "mysql",
      "port": 3306
    }
```

1.39.3 Using ConfigMaps

As Environment Variables:

```
apiVersion: v1
kind: Pod
metadata:
  name: app
spec:
  containers:
  - name: app
    image: myapp:v1
    env:
    - name: ENVIRONMENT
      valueFrom:
        configMapKeyRef:
          name: app-config
          key: ENV
  envFrom:
```

```
- configMapRef:  
  name: app-config
```

As Volume:

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: app  
spec:  
  containers:  
    - name: app  
      image: myapp:v1  
      volumeMounts:  
        - name: config-volume  
          mountPath: /etc/config  
  volumes:  
    - name: config-volume  
      configMap:  
        name: app-config  
        items:  
          - key: config.json  
            path: app-config.json
```

1.40 Secrets

1.40.1 Creating Secrets

```
# From literal values  
kubectl create secret generic db-secret --from-literal=username=admin --from-literal=password=secret123  
  
# From file  
kubectl create secret generic tls-secret --from-file=tls.crt --from-file=tls.key  
  
# Docker registry secret  
kubectl create secret docker-registry regcred --docker-server=https://index.docker.io/v1/ --docker-username=user --docker-password=pass --docker-email=user@example.com  
  
# TLS secret  
kubectl create secret tls tls-secret --cert=tls.crt --key=tls.key
```

1.40.2 Secret YAML

```
apiVersion: v1  
kind: Secret  
metadata:
```

```

    name: db-secret
type: Opaque
data:
  username: YWRtaW4=      # base64 encoded
  password: c2VjcmV0MTIz # base64 encoded
---
# Using stringData (auto-encoded)
apiVersion: v1
kind: Secret
metadata:
  name: db-secret
type: Opaque
stringData:
  username: admin
  password: secret123

```

1.40.3 Secret Types

| Type | Description |
|--------------------------------|-----------------------------|
| Opaque | Generic secret (default) |
| kubernetes.io/dockerconfigjson | Docker registry credentials |
| kubernetes.io/tls | TLS certificate and key |
| kubernetes.io/basic-auth | Basic authentication |
| kubernetes.io/ssh-auth | SSH authentication |

1.40.4 Using Secrets

As Environment Variables:

```

apiVersion: v1
kind: Pod
metadata:
  name: app
spec:
  containers:
  - name: app
    image: myapp:v1
    env:
    - name: DB_PASSWORD
      valueFrom:
        secretKeyRef:
          name: db-secret
          key: password
    envFrom:
    - secretRef:
        name: db-secret

```

As Volume:

```
apiVersion: v1
kind: Pod
metadata:
  name: app
spec:
  containers:
    - name: app
      image: myapp:v1
      volumeMounts:
        - name: secret-volume
          mountPath: /etc/secrets
          readOnly: true
  volumes:
    - name: secret-volume
      secret:
        secretName: db-secret
        defaultMode: 0400
```

1.41 ServiceAccounts

1.41.1 Creating ServiceAccounts

```
kubectl create serviceaccount my-sa
```

1.41.2 ServiceAccount YAML

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: my-sa
automountServiceAccountToken: false
```

1.41.3 Using ServiceAccounts

```
apiVersion: v1
kind: Pod
metadata:
  name: app
spec:
  serviceAccountName: my-sa
  automountServiceAccountToken: true
  containers:
    - name: app
      image: myapp:v1
```

1.42 Security Context

1.42.1 Pod-level Security Context

```
apiVersion: v1
kind: Pod
metadata:
  name: secure-pod
spec:
  securityContext:
    runAsUser: 1000
    runAsGroup: 3000
    fsGroup: 2000
    runAsNonRoot: true
  containers:
  - name: app
    image: myapp:v1
```

1.42.2 Container-level Security Context

```
apiVersion: v1
kind: Pod
metadata:
  name: secure-pod
spec:
  containers:
  - name: app
    image: myapp:v1
    securityContext:
      runAsUser: 1000
      runAsNonRoot: true
      readOnlyRootFilesystem: true
      allowPrivilegeEscalation: false
    capabilities:
      drop:
      - ALL
      add:
      - NET_BIND_SERVICE
```

1.42.3 Security Context Fields

| Field | Level | Description |
|--------------------------|---------------|------------------------------|
| runAsUser | Pod/Container | UID to run as |
| runAsGroup | Pod/Container | GID to run as |
| runAsNonRoot | Pod/Container | Must run as non-root |
| fsGroup | Pod | Group for volumes |
| readOnlyRootFilesystem | Container | Read-only root FS |
| allowPrivilegeEscalation | Container | Prevent privilege escalation |

| Field | Level | Description |
|--------------|-----------|--------------------|
| capabilities | Container | Linux capabilities |

1.43 Resource Requirements

1.43.1 Requests and Limits

```
apiVersion: v1
kind: Pod
metadata:
  name: app
spec:
  containers:
  - name: app
    image: myapp:v1
    resources:
      requests:
        memory: "64Mi"
        cpu: "250m"
      limits:
        memory: "128Mi"
        cpu: "500m"
```

1.43.2 CPU Units

| Value | Description |
|-------|-------------------------------|
| 1 | 1 CPU core |
| 500m | 0.5 CPU core (500 millicores) |
| 100m | 0.1 CPU core |

1.43.3 Memory Units

| Value | Description |
|-------|---------------|
| 128Mi | 128 Mebibytes |
| 1Gi | 1 Gibibyte |
| 256M | 256 Megabytes |

1.43.4 QoS Classes

| Class | Condition |
|------------|--------------------------------------|
| Guaranteed | requests = limits for all containers |
| Burstable | At least one request or limit set |
| BestEffort | No requests or limits set |

1.44 LimitRange

```
apiVersion: v1
kind: LimitRange
metadata:
  name: resource-limits
spec:
  limits:
  - type: Container
    default:
      cpu: "500m"
      memory: "256Mi"
    defaultRequest:
      cpu: "100m"
      memory: "128Mi"
    max:
      cpu: "2"
      memory: "1Gi"
    min:
      cpu: "50m"
      memory: "64Mi"
```

1.45 ResourceQuota

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: compute-quota
spec:
  hard:
    requests.cpu: "4"
    requests.memory: "8Gi"
    limits.cpu: "8"
    limits.memory: "16Gi"
    pods: "10"
    configmaps: "10"
    secrets: "10"
    persistentvolumeclaims: "5"
```

1.46 Admission Controllers

1.46.1 Pod Security Standards

```
apiVersion: v1
kind: Namespace
metadata:
  name: secure-ns
  labels:
    pod-security.kubernetes.io/enforce: restricted
```

```
pod-security.kubernetes.io/audit: restricted  
pod-security.kubernetes.io/warn: restricted
```

| Level | Description |
|------------|-----------------------|
| privileged | Unrestricted |
| baseline | Minimally restrictive |
| restricted | Highly restrictive |

1.47 Key Concepts to Remember

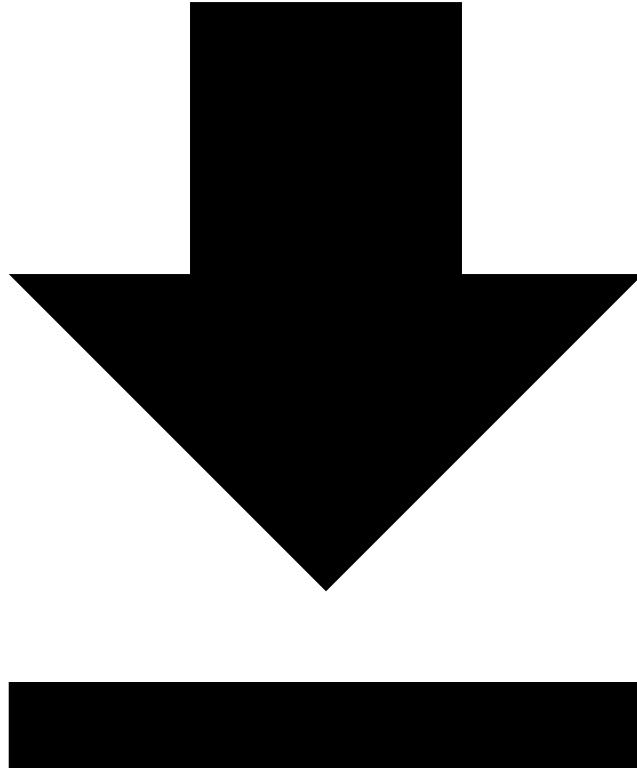
1. **ConfigMaps** - Non-sensitive configuration data
2. **Secrets** - Sensitive data (base64 encoded, not encrypted)
3. **SecurityContext** - Pod/container security settings
4. **Resources** - requests (scheduling) vs limits (enforcement)
5. **ServiceAccounts** - Identity for pods

1.48 Practice Questions

1. How do you create a ConfigMap from a file?
2. What is the difference between `data` and `stringData` in Secrets?
3. How do you mount a Secret as a volume with specific permissions?
4. What QoS class is assigned when requests equal limits?
5. How do you prevent a container from running as root?

[← Previous: Application Observability and Maintenance](#) | [Back to CKAD Overview](#)
| [Next: Services and Networking →](#)

1.49 Services and Networking



[Download PDF Version](#)

This domain covers Kubernetes networking concepts, Services, and Ingress.

1.50 Services

1.50.1 Service Types

| Type | Description |
|--------------|---|
| ClusterIP | Internal cluster IP (default) |
| NodePort | Expose on each node's IP at a static port |
| LoadBalancer | External load balancer (cloud provider) |
| ExternalName | Maps to external DNS name |

1.50.2 ClusterIP Service

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  type: ClusterIP
  selector:
    app: myapp
  ports:
  - port: 80
    targetPort: 8080
    protocol: TCP
```

1.50.3 NodePort Service

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  type: NodePort
  selector:
    app: myapp
  ports:
  - port: 80
    targetPort: 8080
    nodePort: 30080 # Optional: 30000-32767
```

1.50.4 LoadBalancer Service

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  type: LoadBalancer
  selector:
    app: myapp
  ports:
  - port: 80
    targetPort: 8080
```

1.50.5 ExternalName Service

```
apiVersion: v1
kind: Service
metadata:
  name: external-db
spec:
```

```
type: ExternalName
externalName: db.example.com
```

1.50.6 Headless Service

```
apiVersion: v1
kind: Service
metadata:
  name: headless-service
spec:
  clusterIP: None
  selector:
    app: myapp
  ports:
  - port: 80
    targetPort: 8080
```

1.50.7 Creating Services Imperatively

```
# Expose deployment
kubectl expose deployment nginx --port=80 --target-port=8080 --
  type=ClusterIP

# Expose pod
kubectl expose pod nginx --port=80 --target-port=8080

# Create service without selector
kubectl create service clusterip my-svc --tcp=80:8080

# Generate YAML
kubectl expose deployment nginx --port=80 --dry-run=client -o
  yaml > svc.yaml
```

1.51 DNS in Kubernetes

1.51.1 Service DNS

<service-name>.<namespace>.svc.cluster.local

Examples:

- my-service.default.svc.cluster.local
- my-service.default.svc
- my-service.default
- my-service (within same namespace)

1.51.2 Pod DNS

<pod-ip-dashed>.<namespace>.pod.cluster.local

Example:

- 10-244-0-5.default.pod.cluster.local

1.51.3 DNS Resolution Example

```
apiVersion: v1
kind: Pod
metadata:
  name: dns-test
spec:
  containers:
  - name: test
    image: busybox:1.36
    command: ['sleep', '3600']
---
# Test DNS
kubectl exec dns-test -- nslookup my-service
kubectl exec dns-test -- nslookup my-
  service.default.svc.cluster.local
```

1.52 Network Policies

1.52.1 Default Behavior

By default, all pods can communicate with all other pods.

1.52.2 Deny All Ingress

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: deny-all-ingress
  namespace: default
spec:
  podSelector: {}
  policyTypes:
  - Ingress
```

1.52.3 Deny All Egress

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: deny-all-egress
  namespace: default
spec:
  podSelector: {}
```

```
policyTypes:
```

- Egress

1.52.4 Allow Specific Ingress

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-frontend
  namespace: default
spec:
  podSelector:
    matchLabels:
      app: backend
  policyTypes:
    - Ingress
  ingress:
    - from:
        - podSelector:
            matchLabels:
              app: frontend
        - namespaceSelector:
            matchLabels:
              name: production
  ports:
    - protocol: TCP
      port: 8080
```

1.52.5 Allow Egress to Specific Pods

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-db-egress
  namespace: default
spec:
  podSelector:
    matchLabels:
      app: backend
  policyTypes:
    - Egress
  egress:
    - to:
        - podSelector:
            matchLabels:
              app: database
  ports:
    - protocol: TCP
      port: 5432
    - to: # Allow DNS
      - namespaceSelector: {}
```

```

    podSelector:
      matchLabels:
        k8s-app: kube-dns
  ports:
  - protocol: UDP
    port: 53

```

1.52.6 Network Policy Selectors

| Selector | Description |
|-------------------|-----------------------------|
| podSelector | Select pods by labels |
| namespaceSelector | Select namespaces by labels |
| ipBlock | Select by IP CIDR |

```

# IP Block example
ingress:
- from:
  - ipBlock:
    cidr: 10.0.0.0/8
    except:
    - 10.0.1.0/24

```

1.53 Ingress

1.53.1 Ingress Resource

```

apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: my-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  ingressClassName: nginx
  rules:
  - host: myapp.example.com
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: myapp-service
            port:
              number: 80

```

1.53.2 Multiple Hosts

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: multi-host-ingress
spec:
  ingressClassName: nginx
  rules:
  - host: app1.example.com
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: app1-service
            port:
              number: 80
  - host: app2.example.com
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: app2-service
            port:
              number: 80
```

1.53.3 Path-based Routing

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: path-ingress
spec:
  ingressClassName: nginx
  rules:
  - host: myapp.example.com
    http:
      paths:
      - path: /api
        pathType: Prefix
        backend:
          service:
            name: api-service
            port:
              number: 80
      - path: /web
        pathType: Prefix
        backend:
```

```

service:
  name: web-service
  port:
    number: 80

```

1.53.4 TLS Ingress

```

apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: tls-ingress
spec:
  ingressClassName: nginx
  tls:
  - hosts:
    - myapp.example.com
    secretName: tls-secret
  rules:
  - host: myapp.example.com
    http:
      paths:
      - path: /
        pathType: Prefix
      backend:
        service:
          name: myapp-service
          port:
            number: 80

```

1.53.5 Path Types

| Type | Description |
|------------------------|----------------------------------|
| Exact | Exact match of the URL path |
| Prefix | Matches based on URL path prefix |
| ImplementationSpecific | Depends on IngressClass |

1.53.6 Creating Ingress Imperatively

```

# Create ingress
kubectl create ingress my-ingress \
--rule="myapp.example.com/=myapp-service:80" \
--class=nginx

# With TLS
kubectl create ingress my-ingress \
--rule="myapp.example.com/=myapp-service:80,tls=tls-secret" \
--class=nginx

```

1.54 Port Forwarding

```
# Forward pod port
kubectl port-forward pod/nginx 8080:80

# Forward service port
kubectl port-forward svc/nginx 8080:80

# Forward deployment port
kubectl port-forward deployment/nginx 8080:80

# Listen on all interfaces
kubectl port-forward --address 0.0.0.0 pod/nginx 8080:80
```

1.55 Key Concepts to Remember

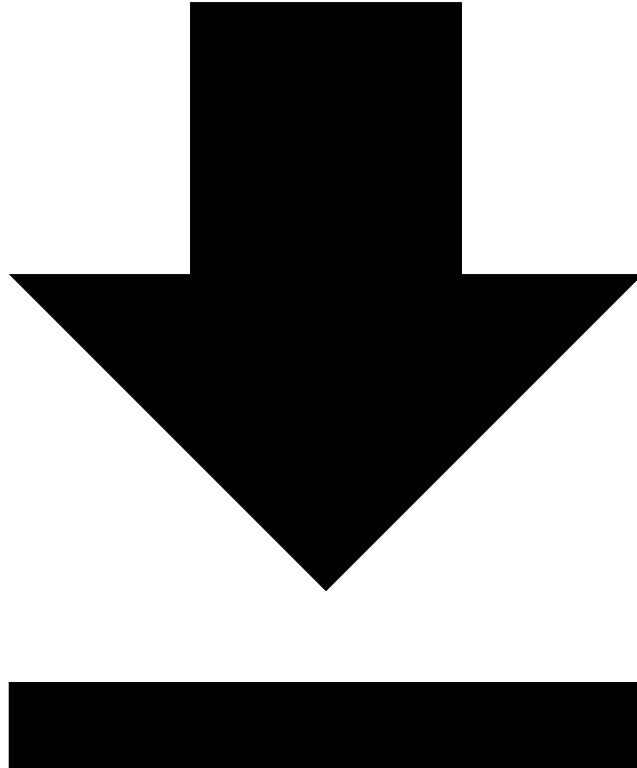
1. **ClusterIP** - Default, internal only
2. **NodePort** - External access via node IP:port
3. **LoadBalancer** - Cloud provider load balancer
4. **Network Policies** - Default allow, explicit deny
5. **Ingress** - HTTP/HTTPS routing, requires controller

1.56 Practice Questions

1. What is the default Service type in Kubernetes?
2. How do you create a Service that exposes a deployment on port 80?
3. What happens to traffic if no Network Policy exists?
4. How do you route traffic based on URL path using Ingress?
5. What is the DNS name format for a Service?

[← Previous: Application Environment, Configuration and Security](#) | [Back to CKAD Overview](#) | [Next: Sample Practice Questions →](#)

1.57 Sample Practice Questions



[Download PDF Version](#)

Disclaimer: These are sample practice questions created for study purposes only. They are NOT actual exam questions and are designed to help you test your understanding of CKAD concepts. Real exam questions may differ in format and content.

1.58 Practice Resources

Before attempting these questions, we highly recommend practicing on:

- [**Killercoda CKAD Scenarios**](#) ★ Free hands-on practice environments
- [**killer.sh CKAD Simulator**](#) - Included with exam registration

1.59 Instructions

- The CKAD exam is **performance-based** (hands-on), not multiple choice
- Practice these scenarios in a real Kubernetes cluster

- Time yourself - aim for efficiency
 - Use imperative commands when possible to save time
-

1.60 Section 1: Application Design and Build (20%)

1.60.1 Question 1.1 - Create a Job

Create a Job named pi-calculator that:

- Uses the image perl:5.34
- Runs the command perl -Mbignum=bpi -wle 'print bpi(2000)'
- Completes 3 times successfully
- Runs 2 pods in parallel
- Has a backoff limit of 4

Show Solution

```
# Generate base YAML
kubectl create job pi-calculator --image=perl:5.34 --dry-
    run=client -o yaml -- perl -Mbignum=bpi -wle 'print
    bpi(2000)' > job.yaml

# Edit to add completions, parallelism, backoffLimit

apiVersion: batch/v1
kind: Job
metadata:
  name: pi-calculator
spec:
  completions: 3
  parallelism: 2
  backoffLimit: 4
  template:
    spec:
      containers:
        - name: pi-calculator
          image: perl:5.34
          command: ["perl", "-Mbignum=bpi", "-wle", "print
                    bpi(2000)"]
          restartPolicy: Never

kubectl apply -f job.yaml
```

1.60.2 Question 1.2 - Create a CronJob

Create a CronJob named backup-job that:

- Runs every day at 2:30 AM
- Uses image busybox:1.36
- Runs command echo "Backup completed at \$(date)"
- Keeps 3 successful job history
- Keeps 1 failed job history

Show Solution

```

kubectl create cronjob backup-job --image=busybox:1.36 --
    schedule="30 2 * * *" -- /bin/sh -c 'echo "Backup
completed at $(date)"'

# Or with YAML for history limits:

apiVersion: batch/v1
kind: CronJob
metadata:
  name: backup-job
spec:
  schedule: "30 2 * * *"
  successfulJobsHistoryLimit: 3
  failedJobsHistoryLimit: 1
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: backup
              image: busybox:1.36
              command: ["/bin/sh", "-c", "echo
\"Backup completed at $(date)\""]
              restartPolicy: OnFailure

```

1.60.3 Question 1.3 - Multi-Container Pod

Create a Pod named `app-with-sidecar` with:
- Main container: `nginx:1.21`
named `main-app` - Sidecar container: `busybox:1.36` named `log-agent` that runs `tail -f /var/log/nginx/access.log` - Both containers share a volume mounted at `/var/log/nginx`

Show Solution

```

apiVersion: v1
kind: Pod
metadata:
  name: app-with-sidecar
spec:
  containers:
    - name: main-app
      image: nginx:1.21
      volumeMounts:
        - name: logs
          mountPath: /var/log/nginx
    - name: log-agent
      image: busybox:1.36
      command: ["tail", "-f", "/var/log/nginx/access.log"]
      volumeMounts:
        - name: logs
          mountPath: /var/log/nginx
  volumes:

```

```
- name: logs  
  emptyDir: {}
```

1.61 Section 2: Application Deployment (20%)

1.61.1 Question 2.1 - Create and Scale Deployment

1. Create a Deployment named web-app with image nginx:1.21 and 3 replicas
2. Update the image to nginx:1.22
3. Check the rollout status
4. Rollback to the previous version

Show Solution

```
# Create deployment  
kubectl create deployment web-app --image=nginx:1.21 --replicas=3  
  
# Update image  
kubectl set image deployment/web-app nginx=nginx:1.22  
  
# Check rollout status  
kubectl rollout status deployment/web-app  
  
# View history  
kubectl rollout history deployment/web-app  
  
# Rollback  
kubectl rollout undo deployment/web-app
```

1.61.2 Question 2.2 - Deployment Strategy

Create a Deployment named rolling-app with:
- Image: nginx:1.21 - 4 replicas
- Rolling update strategy with maxSurge=1 and maxUnavailable=1

Show Solution

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: rolling-app  
spec:  
  replicas: 4  
  strategy:  
    type: RollingUpdate  
    rollingUpdate:  
      maxSurge: 1  
      maxUnavailable: 1  
  selector:  
    matchLabels:  
      app: rolling-app
```

```
template:
  metadata:
    labels:
      app: rolling-app
  spec:
    containers:
      - name: nginx
        image: nginx:1.21
```

1.61.3 Question 2.3 - Helm Operations

1. Add the bitnami repository
2. Search for nginx chart
3. Install nginx chart with release name my-nginx
4. List all releases
5. Uninstall the release

Show Solution

```
# Add repository
helm repo add bitnami https://charts.bitnami.com/bitnami
helm repo update

# Search
helm search repo nginx

# Install
helm install my-nginx bitnami/nginx

# List releases
helm list

# Uninstall
helm uninstall my-nginx
```

1.62 Section 3: Application Observability and Maintenance (15%)

1.62.1 Question 3.1 - Configure Probes

Create a Pod named health-check-pod with: - Image: nginx:1.21 - Liveness probe: HTTP GET on path / port 80, initial delay 10s, period 5s - Readiness probe: HTTP GET on path / port 80, initial delay 5s, period 3s

Show Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: health-check-pod
```

```

spec:
  containers:
  - name: nginx
    image: nginx:1.21
    ports:
    - containerPort: 80
    livenessProbe:
      httpGet:
        path: /
        port: 80
    initialDelaySeconds: 10
    periodSeconds: 5
    readinessProbe:
      httpGet:
        path: /
        port: 80
    initialDelaySeconds: 5
    periodSeconds: 3

```

1.62.2 Question 3.2 - Debugging

A Pod named broken-pod is not running correctly. Debug and fix it.

```

apiVersion: v1
kind: Pod
metadata:
  name: broken-pod
spec:
  containers:
  - name: app
    image: nginx:latest
    command: ["nginx", "-g", "daemon off;"]
    resources:
      limits:
        memory: "10Mi"

```

Show Solution

```

# Check pod status
kubectl get pod broken-pod

# Check events and details
kubectl describe pod broken-pod

# Check logs
kubectl logs broken-pod

# The issue is likely OOMKilled due to low memory limit
# Fix by increasing memory limit:

apiVersion: v1
kind: Pod
metadata:

```

```
  name: broken-pod
spec:
  containers:
    - name: app
      image: nginx:latest
      resources:
        limits:
          memory: "128Mi"
```

1.62.3 Question 3.3 - View Logs

1. View logs of pod nginx in namespace web
2. View logs of the previous container instance
3. Follow logs in real-time
4. View last 50 lines

Show Solution

```
# View logs
kubectl logs nginx -n web

# Previous container
kubectl logs nginx -n web --previous

# Follow logs
kubectl logs -f nginx -n web

# Last 50 lines
kubectl logs nginx -n web --tail=50
```

1.63 Section 4: Application Environment, Configuration and Security (25%)

1.63.1 Question 4.1 - ConfigMap and Secret

1. Create a ConfigMap named app-config with:
 - APP_ENV=production
 - LOG_LEVEL=info
2. Create a Secret named db-secret with:
 - DB_USER=admin
 - DB_PASS=secret123
3. Create a Pod that uses both as environment variables

Show Solution

```
# Create ConfigMap
kubectl create configmap app-config --from-
  literal=APP_ENV=production --from-literal=LOG_LEVEL=info
```

```

# Create Secret
kubectl create secret generic db-secret --from-
    literal=DB_USER=admin --from-literal=DB_PASS=secret123

apiVersion: v1
kind: Pod
metadata:
  name: app-pod
spec:
  containers:
    - name: app
      image: nginx:1.21
      envFrom:
        - configMapRef:
            name: app-config
        - secretRef:
            name: db-secret

```

1.63.2 Question 4.2 - Security Context

Create a Pod named `secure-pod` that:

- Runs as user ID 1000
- Runs as group ID 3000
- Has a read-only root filesystem
- Cannot escalate privileges

Show Solution

```

apiVersion: v1
kind: Pod
metadata:
  name: secure-pod
spec:
  securityContext:
    runAsUser: 1000
    runAsGroup: 3000
  containers:
    - name: app
      image: nginx:1.21
      securityContext:
        readOnlyRootFilesystem: true
        allowPrivilegeEscalation: false

```

1.63.3 Question 4.3 - Resource Limits

Create a Pod named `resource-pod` with:

- Image: `nginx:1.21`
- CPU request: 100m, limit: 200m
- Memory request: 64Mi, limit: 128Mi

Show Solution

```

kubectl run resource-pod --image=nginx:1.21 --dry-run=client -o
    yaml > pod.yaml
# Edit to add resources

apiVersion: v1
kind: Pod

```

```

metadata:
  name: resource-pod
spec:
  containers:
    - name: nginx
      image: nginx:1.21
      resources:
        requests:
          cpu: "100m"
          memory: "64Mi"
        limits:
          cpu: "200m"
          memory: "128Mi"

```

1.64 Section 5: Services and Networking (20%)

1.64.1 Question 5.1 - Create Services

1. Create a Deployment named web with image nginx:1.21 and 3 replicas
2. Expose it as a ClusterIP Service on port 80
3. Expose it as a NodePort Service on port 30080

Show Solution

```

# Create deployment
kubectl create deployment web --image=nginx:1.21 --replicas=3

# ClusterIP service
kubectl expose deployment web --port=80 --target-port=80 --
  name=web-clusterip

# NodePort service
kubectl expose deployment web --port=80 --target-port=80 --
  type=NodePort --name=web-nodeport

# Or specify nodePort:
kubectl create service nodeport web-nodeport --tcp=80:80 --node-
  port=30080

```

1.64.2 Question 5.2 - Network Policy

Create a NetworkPolicy named api-policy in namespace default that:

- Applies to pods with label app=api
- Allows ingress only from pods with label app=frontend
- Allows ingress only on port 8080

Show Solution

```

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:

```

```

name: api-policy
namespace: default
spec:
  podSelector:
    matchLabels:
      app: api
  policyTypes:
    - Ingress
  ingress:
    - from:
        - podSelector:
            matchLabels:
              app: frontend
  ports:
    - protocol: TCP
      port: 8080

```

1.64.3 Question 5.3 - Ingress

Create an Ingress named web-ingress that:

- Routes `app.example.com/api` to service `api-service` port 80
- Routes `app.example.com/web` to service `web-service` port 80
- Uses ingress class `nginx`

Show Solution

```

apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: web-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  ingressClassName: nginx
  rules:
    - host: app.example.com
      http:
        paths:
          - path: /api
            pathType: Prefix
            backend:
              service:
                name: api-service
                port:
                  number: 80
          - path: /web
            pathType: Prefix
            backend:
              service:
                name: web-service
                port:
                  number: 80

```

1.65 Exam Tips

1. **Use aliases:** alias k=kubectl
2. **Enable auto-completion:** source <(kubectl completion bash)
3. **Use --dry-run=client -o yaml** to generate YAML templates
4. **Bookmark important docs** before the exam
5. **Practice on Killercoda** for free hands-on scenarios
6. **Time management:** Don't spend too long on any single question
7. **Use imperative commands** when possible to save time

1.66 Additional Practice

- [Killercoda CKAD Scenarios](#) - Free interactive scenarios
 - [killer.sh](#) - Exam simulator (included with registration)
 - [Kubernetes Documentation](#) - Allowed during exam
-

[← Back to CKAD Overview](#)
