

## 1 CKS

- 1.1 Table of Contents
- 1.2 Overview
- 1.3 Exam Overview
- 1.4 Exam Domains & Weights
- 1.5 Prerequisites
- 1.6 Study Resources
  - 1.6.1 Official Resources
  - 1.6.2 Recommended Courses
  - 1.6.3 Practice Resources
- 1.7 Quick Navigation
- 1.8 Exam Environment
  - 1.8.1 Allowed Resources During Exam
- 1.9 Exam Tips
- 1.10 Security Tools Overview
- 1.11 Useful Commands
- 1.12 Registration
- 1.13 Cluster Setup
- 1.14 Network Policies
  - 1.14.1 Default Deny All Ingress
  - 1.14.2 Default Deny All Egress
  - 1.14.3 Default Deny All Traffic
  - 1.14.4 Allow Specific Ingress
  - 1.14.5 Allow DNS Egress
- 1.15 CIS Kubernetes Benchmark
  - 1.15.1 kube-bench
  - 1.15.2 Common CIS Findings
- 1.16 Ingress Security
  - 1.16.1 TLS Ingress
  - 1.16.2 Create TLS Secret
- 1.17 Cluster Component Security
  - 1.17.1 Secure API Server
  - 1.17.2 Secure etcd
  - 1.17.3 Secure kubelet
- 1.18 Encryption at Rest
  - 1.18.1 Encryption Configuration
  - 1.18.2 Enable Encryption
  - 1.18.3 Verify Encryption
- 1.19 GUI Element Security
  - 1.19.1 Disable Kubernetes Dashboard (if not needed)
  - 1.19.2 Secure Dashboard Access
- 1.20 Key Concepts to Remember
- 1.21 Practice Questions
- 1.22 Cluster Hardening
- 1.23 RBAC (Role-Based Access Control)
  - 1.23.1 Role (Namespace-scoped)
  - 1.23.2 ClusterRole (Cluster-scoped)
  - 1.23.3 RoleBinding
  - 1.23.4 ClusterRoleBinding
  - 1.23.5 RBAC Commands
  - 1.23.6 Least Privilege Principle

- 1.24 ServiceAccount Security
  - 1.24.1 Create ServiceAccount
  - 1.24.2 Disable Token Auto-mount
  - 1.24.3 ServiceAccount Token Projection
- 1.25 Restrict API Access
  - 1.25.1 Disable Anonymous Authentication
  - 1.25.2 Enable RBAC Authorization
  - 1.25.3 Restrict kubelet API
- 1.26 Upgrade Kubernetes
  - 1.26.1 Upgrade Control Plane
  - 1.26.2 Upgrade Worker Nodes
- 1.27 Restrict Access to Kubernetes API
  - 1.27.1 API Server Flags
  - 1.27.2 Network-level Restrictions
- 1.28 Admission Controllers
  - 1.28.1 Important Security Admission Controllers
  - 1.28.2 Enable Admission Controllers
- 1.29 Key Concepts to Remember
- 1.30 Practice Questions
- 1.31 System Hardening
- 1.32 AppArmor
  - 1.32.1 AppArmor Basics
  - 1.32.2 AppArmor Profile Example
  - 1.32.3 Use AppArmor in Pod
  - 1.32.4 AppArmor Profile Modes
- 1.33 seccomp (Secure Computing Mode)
  - 1.33.1 seccomp Profile Location
  - 1.33.2 seccomp Profile Example
  - 1.33.3 Use seccomp in Pod
  - 1.33.4 seccomp Profile Types
  - 1.33.5 RuntimeDefault seccomp
- 1.34 Minimize Host OS Footprint
  - 1.34.1 Reduce Attack Surface
  - 1.34.2 Secure SSH
  - 1.34.3 File System Security
- 1.35 Limit Node Access
  - 1.35.1 Restrict SSH Access
  - 1.35.2 Use Network Policies for Node Access
- 1.36 Linux Capabilities
  - 1.36.1 Drop All Capabilities
  - 1.36.2 Add Specific Capabilities
  - 1.36.3 Common Capabilities
- 1.37 Kernel Hardening
  - 1.37.1 sysctl Settings
- 1.38 IAM Roles (Cloud)
  - 1.38.1 AWS IAM for Nodes
  - 1.38.2 Pod Identity (IRSA)
- 1.39 Key Concepts to Remember
- 1.40 Practice Questions
- 1.41 Minimize Microservice Vulnerabilities
- 1.42 Pod Security Standards
  - 1.42.1 Pod Security Admission
  - 1.42.2 Security Levels
  - 1.42.3 Restricted Pod Example

- 1.43 Security Context
  - 1.43.1 Pod-level Security Context
  - 1.43.2 Container-level Security Context
  - 1.43.3 Security Context Fields
- 1.44 OPA Gatekeeper
  - 1.44.1 Install Gatekeeper
  - 1.44.2 Constraint Template
  - 1.44.3 Constraint
  - 1.44.4 Deny Privileged Containers
- 1.45 Secrets Management
  - 1.45.1 Create Secrets
  - 1.45.2 Use Secrets as Environment Variables
  - 1.45.3 Use Secrets as Volumes
  - 1.45.4 Encrypt Secrets at Rest
- 1.46 Container Runtime Sandboxing
  - 1.46.1 gVisor (runsc)
  - 1.46.2 Kata Containers
- 1.47 mTLS with Service Mesh
  - 1.47.1 Istio Strict mTLS
- 1.48 Key Concepts to Remember
- 1.49 Practice Questions
- 1.50 Supply Chain Security
- 1.51 Image Vulnerability Scanning
  - 1.51.1 Trivy
  - 1.51.2 Trivy Output Example
  - 1.51.3 Integrate Trivy in CI/CD
- 1.52 Image Signing and Verification
  - 1.52.1 Cosign
  - 1.52.2 Image Policy Webhook
- 1.53 Minimize Base Image Footprint
  - 1.53.1 Use Minimal Base Images
  - 1.53.2 Multi-stage Builds
  - 1.53.3 Image Best Practices
- 1.54 Secure Image Registries
  - 1.54.1 Private Registry Authentication
  - 1.54.2 Use ImagePullSecrets
  - 1.54.3 ServiceAccount with ImagePullSecrets
- 1.55 Allowlist Registries
  - 1.55.1 OPA Gatekeeper Policy
- 1.56 Static Analysis
  - 1.56.1 Kubesec
  - 1.56.2 Kubesec Output
  - 1.56.3 Conftest (OPA for Config Files)
  - 1.56.4 Example Conftest Policy
- 1.57 Dockerfile Security
  - 1.57.1 Dockerfile Best Practices
  - 1.57.2 Hadolint (Dockerfile Linter)
- 1.58 Key Concepts to Remember
- 1.59 Practice Questions
- 1.60 Monitoring, Logging and Runtime Security
- 1.61 Overview
- 1.62 Falco - Runtime Threat Detection
  - 1.62.1 Install Falco
  - 1.62.2 Falco Architecture

- 1.62.3 Falco Rules Structure
- 1.62.4 Common Falco Rules
- 1.62.5 Custom Falco Rules
- 1.62.6 Falco Commands
- 1.63 Kubernetes Audit Logging
  - 1.63.1 Audit Stages
  - 1.63.2 Audit Levels
  - 1.63.3 Comprehensive Audit Policy
  - 1.63.4 Enable Audit Logging on API Server
  - 1.63.5 Analyze Audit Logs
  - 1.63.6 Audit Log Entry Structure
- 1.64 Container Immutability
  - 1.64.1 Read-Only Root Filesystem
  - 1.64.2 Enforce Immutability with Pod Security Standards
- 1.65 Sysdig and Other Tools
  - 1.65.1 Sysdig for Troubleshooting
- 1.66 Key Concepts Summary
- 1.67 Practice Exercises
- 1.68 Sample Practice Questions
- 1.69 Practice Resources
- 1.70 Section 1: Cluster Setup (10%)
  - 1.70.1 Question 1.1 - Network Policy
  - 1.70.2 Question 1.2 - CIS Benchmark
- 1.71 Section 2: Cluster Hardening (15%)
  - 1.71.1 Question 2.1 - RBAC
  - 1.71.2 Question 2.2 - ServiceAccount
- 1.72 Section 3: System Hardening (15%)
  - 1.72.1 Question 3.1 - AppArmor
  - 1.72.2 Question 3.2 - seccomp
- 1.73 Section 4: Minimize Microservice Vulnerabilities (20%)
  - 1.73.1 Question 4.1 - Security Context
  - 1.73.2 Question 4.2 - Pod Security Standards
- 1.74 Section 5: Supply Chain Security (20%)
  - 1.74.1 Question 5.1 - Image Scanning
  - 1.74.2 Question 5.2 - Allowed Registries
- 1.75 Section 6: Monitoring, Logging and Runtime Security (20%)
  - 1.75.1 Question 6.1 - Audit Logging
  - 1.75.2 Question 6.2 - Falco
  - 1.75.3 Question 6.3 - Container Immutability
  - 1.75.4 Question 6.4 - Analyze Audit Logs
- 1.76 Section 7: Additional Practice Questions
  - 1.76.1 Question 7.1 - Ingress TLS
  - 1.76.2 Question 7.2 - RuntimeClass with gVisor
  - 1.76.3 Question 7.3 - Encrypt etcd Data
  - 1.76.4 Question 7.4 - Verify Image Signature
  - 1.76.5 Question 7.5 - Restrict Syscalls with Seccomp
  - 1.76.6 Question 7.6 - Network Policy for Database
  - 1.76.7 Question 7.7 - Admission Controller Webhook
  - 1.76.8 Question 7.8 - Investigate Compromised Pod
- 1.77 Exam Tips
- 1.78 Quick Reference Commands

# 1 CKS

Generated on: 2026-01-13 15:04:35 Version: 1.0

---

## 1.1 Table of Contents

1. [Overview](#)
  2. [Cluster Setup](#)
  3. [Cluster Hardening](#)
  4. [System Hardening](#)
  5. [Minimize Microservice Vulnerabilities](#)
  6. [Supply Chain Security](#)
  7. [Monitoring, Logging and Runtime Security](#)
  8. [Sample Practice Questions](#)
- 

## 1.2 Overview



The Certified Kubernetes Security Specialist (CKS) exam certifies that candidates have the skills, knowledge, and competency to perform a broad range of best practices for securing container-based applications and Kubernetes platforms during build, deployment, and runtime.

## 1.3 Exam Overview

| Detail                        | Information                       |
|-------------------------------|-----------------------------------|
| <b>Exam Format</b>            | Performance-based (hands-on)      |
| <b>Number of Questions</b>    | 15-20                             |
| <b>Duration</b>               | 2 hours                           |
| <b>Passing Score</b>          | 67%                               |
| <b>Certification Validity</b> | 2 years                           |
| <b>Cost</b>                   | \$395 USD                         |
| <b>Retake Policy</b>          | 1 free retake                     |
| <b>Kubernetes Version</b>     | 1.30                              |
| <b>Prerequisite</b>           | Must hold valid CKA certification |

## 1.4 Exam Domains & Weights

| Domain   | Weight |
|--|--------|
| <a href="#">Cluster Setup</a>                            | 10%    |
| <a href="#">Cluster Hardening</a>                        | 15%    |
| <a href="#">System Hardening</a>                         | 15%    |
| <a href="#">Minimize Microservice Vulnerabilities</a>    | 20%    |
| <a href="#">Supply Chain Security</a>                    | 20%    |
| <a href="#">Monitoring, Logging and Runtime Security</a> | 20%    |

## 1.5 Prerequisites

- **Valid CKA certification** (required)
- Strong Kubernetes administration skills
- Understanding of Linux security concepts
- Familiarity with container security

## 1.6 Study Resources

### 1.6.1 Official Resources

- [CKS Exam Curriculum](#)
- [Kubernetes Security Documentation](#)
- [Kubernetes Security Best Practices](#)

### 1.6.2 Recommended Courses

- [Kubernetes Security Essentials \(LFS260\)](#)
- [CKS with Practice Tests - Udemy](#)

### 1.6.3 Practice Resources

- [Killercoda CKS Scenarios](#) ★ Highly Recommended
- [killer.sh CKS Simulator](#) - Included with exam registration

## 1.7 Quick Navigation

- [01 - Cluster Setup](#)
- [02 - Cluster Hardening](#)
- [03 - System Hardening](#)
- [04 - Minimize Microservice Vulnerabilities](#)
- [05 - Supply Chain Security](#)

- [06 - Monitoring, Logging and Runtime Security](#)
- [Sample Practice Questions](#)

## 1.8 Exam Environment

The CKS exam provides:

- Access to multiple Kubernetes clusters
- `kubectl` with auto-completion enabled
- Access to Kubernetes documentation ([kubernetes.io](https://kubernetes.io))
- A Linux terminal environment
- Root access via `sudo`
- Security tools pre-installed

### 1.8.1 Allowed Resources During Exam

- [kubernetes.io/docs](https://kubernetes.io/docs)
- [kubernetes.io/blog](https://kubernetes.io/blog)
- [trivy.dev/docs](https://trivy.dev/docs)
- [falco.org/docs](https://falco.org/docs)
- [aquasecurity.github.io/trivy](https://aquasecurity.github.io/trivy)

## 1.9 Exam Tips

1. **Master security contexts** - `runAsUser`, `runAsNonRoot`, capabilities
2. **Know Network Policies** - Default deny, allow specific traffic
3. **Understand RBAC deeply** - Least privilege principle
4. **Practice with security tools** - Trivy, Falco, AppArmor, seccomp
5. **Know Pod Security Standards** - privileged, baseline, restricted
6. **Practice image scanning** - Trivy for vulnerability detection
7. **Understand audit logging** - Configure and analyze audit logs
8. **Practice on [Killercoda](#)** - Free hands-on scenarios

## 1.10 Security Tools Overview

| Tool                  | Purpose                                  |
|-----------------------|--|
| <b>Trivy</b>          | Container image vulnerability scanning   |
| <b>Falco</b>          | Runtime security monitoring              |
| <b>AppArmor</b>       | Linux security module for access control |
| <b>seccomp</b>        | System call filtering                    |
| <b>OPA/Gatekeeper</b> | Policy enforcement                       |
| <b>kube-bench</b>     | CIS Kubernetes benchmark                 |

## 1.11 Useful Commands

```
# Set alias
alias k=kubectl

# Enable auto-completion
source <(kubectl completion bash)
complete -o default -F __start_kubectl k

# Check API server audit logs
cat /var/log/kubernetes/audit/audit.log | jq .

# Scan image with Trivy
trivy image nginx:1.21

# Check seccomp profiles
ls /var/lib/kubelet/seccomp/

# View AppArmor profiles
aa-status

# Check Pod Security Standards
kubectl label namespace default pod-security.kubernetes.io/
    enforce=restricted

# RBAC verification
kubectl auth can-i --list --as system:serviceaccount:default:mysa
kubectl auth can-i create pods --as jane

# Network Policy testing
kubectl exec -it test-pod -- nc -zv target-service 80
```


## 1.12 Registration

[Register for CKS Exam](#)

| **Note:** You must hold a valid CKA certification before taking the CKS exam.

---

## 1.13 Cluster Setup



[Download PDF Version](#)

This domain covers securing Kubernetes cluster components and network security.

## 1.14 Network Policies

### 1.14.1 Default Deny All Ingress

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: default-deny-ingress
  namespace: default
spec:
  podSelector: {}
  policyTypes:
    - Ingress
```

## 1.14.2 Default Deny All Egress

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: default-deny-egress
  namespace: default
spec:
  podSelector: {}
  policyTypes:
    - Egress
```

## 1.14.3 Default Deny All Traffic

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: default-deny-all
  namespace: default
spec:
  podSelector: {}
  policyTypes:
    - Ingress
    - Egress
```

## 1.14.4 Allow Specific Ingress

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-frontend-to-backend
  namespace: default
spec:
  podSelector:
    matchLabels:
      app: backend
  policyTypes:
    - Ingress
  ingress:
    - from:
        - podSelector:
            matchLabels:
              app: frontend
        - namespaceSelector:
            matchLabels:
              name: production
  ports:
    - protocol: TCP
      port: 8080
```

## 1.14.5 Allow DNS Egress

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-dns-egress
  namespace: default
spec:
  podSelector: {}
  policyTypes:
    - Egress
  egress:
    - to:
        - namespaceSelector: {}
          podSelector:
            matchLabels:
              k8s-app: kube-dns
        ports:
          - protocol: UDP
            port: 53
          - protocol: TCP
            port: 53
```

## 1.15 CIS Kubernetes Benchmark

### 1.15.1 kube-bench

Tool to check Kubernetes against CIS benchmarks.

```
# Run kube-bench
kube-bench run --targets=master
kube-bench run --targets=node
kube-bench run --targets=etcd

# Run specific checks
kube-bench run --targets=master --check=1.1.1

# Output as JSON
kube-bench run --json
```

### 1.15.2 Common CIS Findings

| Check | Description                     | Fix   |
|-------|---------------------------------|---|
| 1.1.1 | API server pod spec permissions | chmod 600 /etc/kubernetes/manifests/kube-apiserver.yaml |
| 1.2.1 | Anonymous auth disabled         | --anonymous-auth=false                                  |

| Check | Description                      | Fix   |
|-------|----------------------------------|---|
| 1.2.6 | RBAC enabled                     | --authorization-mode=RBAC                                       |
| 4.1.1 | kubelet service file permissions | chmod 600 /etc/systemd/system/kubelet.service.d/10-kubeadm.conf |

## 1.16 Ingress Security

### 1.16.1 TLS Ingress

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: secure-ingress
  annotations:
    nginx.ingress.kubernetes.io/ssl-redirect: "true"
spec:
  ingressClassName: nginx
  tls:
  - hosts:
    - secure.example.com
      secretName: tls-secret
  rules:
  - host: secure.example.com
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: secure-service
            port:
              number: 443
```

### 1.16.2 Create TLS Secret

```
# Generate self-signed certificate
openssl req -x509 -nodes -days 365 -newkey rsa:2048 \
  -keyout tls.key -out tls.crt -subj "/CN=secure.example.com"

# Create secret
kubectl create secret tls tls-secret --cert=tls.crt --key=tls.key
```

# 1.17 Cluster Component Security

## 1.17.1 Secure API Server

```
# /etc/kubernetes/manifests/kube-apiserver.yaml
spec:
  containers:
    - command:
        - kube-apiserver
        - --anonymous-auth=false
        - --authorization-mode=Node,RBAC
        - --enable-admission-plugins=NodeRestriction,PodSecurityPolicy
        - --audit-log-path=/var/log/kubernetes/audit/audit.log
        - --audit-policy-file=/etc/kubernetes/audit-policy.yaml
        - --encryption-provider-config=/etc/kubernetes/encryption-
          config.yaml
        - --tls-cert-file=/etc/kubernetes/pki/apiserver.crt
        - --tls-private-key-file=/etc/kubernetes/pki/apiserver.key
        - --client-ca-file=/etc/kubernetes/pki/ca.crt
```

## 1.17.2 Secure etcd

```
# /etc/kubernetes/manifests/etcd.yaml
spec:
  containers:
    - command:
        - etcd
        - --cert-file=/etc/kubernetes/pki/etcd/server.crt
        - --key-file=/etc/kubernetes/pki/etcd/server.key
        - --trusted-ca-file=/etc/kubernetes/pki/etcd/ca.crt
        - --client-cert-auth=true
        - --peer-cert-file=/etc/kubernetes/pki/etcd/peer.crt
        - --peer-key-file=/etc/kubernetes/pki/etcd/peer.key
        - --peer-trusted-ca-file=/etc/kubernetes/pki/etcd/ca.crt
        - --peer-client-cert-auth=true
```

## 1.17.3 Secure kubelet

```
# /var/lib/kubelet/config.yaml
authentication:
  anonymous:
    enabled: false
  webhook:
    enabled: true
  x509:
    clientCAFile: /etc/kubernetes/pki/ca.crt
authorization:
  mode: Webhook
readOnlyPort: 0
```

## 1.18 Encryption at Rest

### 1.18.1 Encryption Configuration

```
# /etc/kubernetes/encryption-config.yaml
apiVersion: apiserver.config.k8s.io/v1
kind: EncryptionConfiguration
resources:
  - resources:
    - secrets
  providers:
    - aescbc:
      keys:
        - name: key1
          secret: <base64-encoded-32-byte-key>
    - identity: {}
```

### 1.18.2 Enable Encryption

```
# Generate encryption key
head -c 32 /dev/urandom | base64

# Add to API server
# --encryption-provider-config=/etc/kubernetes/encryption-
# config.yaml

# Encrypt existing secrets
kubectl get secrets --all-namespaces -o json | kubectl replace -f
-
```

### 1.18.3 Verify Encryption

```
# Check if secrets are encrypted in etcd
ETCDCTL_API=3 etcdctl get /registry/secrets/default/my-secret \
--endpoints=https://127.0.0.1:2379 \
--cacert=/etc/kubernetes/pki/etcd/ca.crt \
--cert=/etc/kubernetes/pki/etcd/server.crt \
--key=/etc/kubernetes/pki/etcd/server.key | hexdump -C
```

## 1.19 GUI Element Security

### 1.19.1 Disable Kubernetes Dashboard (if not needed)

```
# Delete dashboard
kubectl delete ns kubernetes-dashboard

# Or restrict access with NetworkPolicy
```

## 1.19.2 Secure Dashboard Access

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: dashboard-policy
  namespace: kubernetes-dashboard
spec:
  podSelector:
    matchLabels:
      k8s-app: kubernetes-dashboard
  policyTypes:
  - Ingress
  ingress:
  - from:
    - ipBlock:
        cidr: 10.0.0.0/8
    ports:
    - protocol: TCP
      port: 443
```

## 1.20 Key Concepts to Remember

1. **Network Policies** - Default deny, allow specific traffic
2. **CIS Benchmarks** - Use kube-bench to verify compliance
3. **TLS everywhere** - Encrypt all cluster communication
4. **Encryption at rest** - Encrypt secrets in etcd
5. **Disable anonymous auth** - Require authentication

## 1.21 Practice Questions


1. How do you create a default deny all network policy?
2. What tool checks Kubernetes against CIS benchmarks?
3. How do you enable encryption at rest for secrets?
4. What flag disables anonymous authentication on the API server?
5. How do you verify secrets are encrypted in etcd?

---

[Back to CKS Overview](#) | [Next: Cluster Hardening →](#)

---

## 1.22 Cluster Hardening



[Download PDF Version](#)

This domain covers hardening Kubernetes cluster components and implementing access controls.

## 1.23 RBAC (Role-Based Access Control)

### 1.23.1 Role (Namespace-scoped)

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  namespace: default
  name: pod-reader
rules:
- apiGroups: [""]
  resources: ["pods"]
  verbs: ["get", "watch", "list"]
- apiGroups: ["]
```

```
resources: ["pods/log"]
verbs: ["get"]
```

### 1.23.2 ClusterRole (Cluster-scoped)

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: secret-reader
rules:
- apiGroups: [""]
  resources: ["secrets"]
  verbs: ["get", "watch", "list"]
- apiGroups: [""]
  resources: ["nodes"]
  verbs: ["get", "list"]
```

### 1.23.3 RoleBinding

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: read-pods
  namespace: default
subjects:
- kind: User
  name: jane
  apiGroup: rbac.authorization.k8s.io
- kind: ServiceAccount
  name: default
  namespace: kube-system
- kind: Group
  name: developers
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role
  name: pod-reader
  apiGroup: rbac.authorization.k8s.io
```

### 1.23.4 ClusterRoleBinding

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: read-secrets-global
subjects:
- kind: Group
  name: managers
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: ClusterRole
```

```
name: secret-reader
apiGroup: rbac.authorization.k8s.io
```

### 1.23.5 RBAC Commands

```
# Create role
kubectl create role pod-reader --verb=get,list,watch --
    resource=pods -n default

# Create clusterrole
kubectl create clusterrole node-reader --verb=get,list --
    resource=nodes

# Create rolebinding
kubectl create rolebinding read-pods --role=pod-reader --
    user=jane -n default

# Create clusterrolebinding
kubectl create clusterrolebinding read-nodes --clusterrole=node-
    reader --group=developers

# Check permissions
kubectl auth can-i create pods --as jane
kubectl auth can-i list nodes --as jane
kubectl auth can-i --list --as jane
kubectl auth can-i --list --as system:serviceaccount:default:mysa
```

### 1.23.6 Least Privilege Principle

```
# Bad: Too permissive
rules:
- apiGroups: ["*"]
  resources: ["*"]
  verbs: ["*"]

# Good: Specific permissions
rules:
- apiGroups: [""]
  resources: ["pods"]
  verbs: ["get", "list"]
- apiGroups: ["apps"]
  resources: ["deployments"]
  verbs: ["get", "list"]
```

## 1.24 ServiceAccount Security

### 1.24.1 Create ServiceAccount

```
apiVersion: v1
kind: ServiceAccount
```

```
metadata:
  name: my-sa
  namespace: default
automountServiceAccountToken: false
```

## 1.24.2 Disable Token Auto-mount

```
apiVersion: v1
kind: Pod
metadata:
  name: secure-pod
spec:
  serviceAccountName: my-sa
  automountServiceAccountToken: false
  containers:
    - name: app
      image: nginx
```

## 1.24.3 ServiceAccount Token Projection

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-with-projected-token
spec:
  serviceAccountName: my-sa
  containers:
    - name: app
      image: nginx
      volumeMounts:
        - name: token
          mountPath: /var/run/secrets/tokens
  volumes:
    - name: token
      projected:
        sources:
          - serviceAccountToken:
              path: token
              expirationSeconds: 3600
              audience: api
```

# 1.25 Restrict API Access

## 1.25.1 Disable Anonymous Authentication

```
# /etc/kubernetes/manifests/kube-apiserver.yaml
spec:
  containers:
    - command:
```

```
- kube-apiserver
- --anonymous-auth=false
```

### 1.25.2 Enable RBAC Authorization

```
# /etc/kubernetes/manifests/kube-apiserver.yaml
spec:
  containers:
    - command:
      - kube-apiserver
      - --authorization-mode=Node,RBAC
```

### 1.25.3 Restrict kubelet API

```
# /var/lib/kubelet/config.yaml
authentication:
  anonymous:
    enabled: false
  webhook:
    enabled: true
authorization:
  mode: Webhook
```

## 1.26 Upgrade Kubernetes

### 1.26.1 Upgrade Control Plane

```
# Check available versions
apt-cache madison kubeadm

# Upgrade kubeadm
apt-mark unhold kubeadm
apt-get update && apt-get install -y kubeadm=1.30.0-1.1
apt-mark hold kubeadm

# Plan upgrade
kubeadm upgrade plan

# Apply upgrade
kubeadm upgrade apply v1.30.0

# Upgrade kubelet and kubectl
apt-mark unhold kubelet kubectl
apt-get update && apt-get install -y kubelet=1.30.0-1.1
                           kubectl=1.30.0-1.1
apt-mark hold kubelet kubectl

systemctl daemon-reload
systemctl restart kubelet
```

## 1.26.2 Upgrade Worker Nodes

```
# Drain node
kubectl drain <node-name> --ignore-daemonsets --delete-emptydir-
    data

# On worker node
apt-mark unhold kubeadm
apt-get update && apt-get install -y kubeadm=1.30.0-1.1
apt-mark hold kubeadm

kubeadm upgrade node

apt-mark unhold kubelet kubectl
apt-get update && apt-get install -y kubelet=1.30.0-1.1
    kubectl=1.30.0-1.1
apt-mark hold kubelet kubectl

systemctl daemon-reload
systemctl restart kubelet

# Uncordon node
kubectl uncordon <node-name>
```

## 1.27 Restrict Access to Kubernetes API

### 1.27.1 API Server Flags

```
# /etc/kubernetes/manifests/kube-apiserver.yaml
spec:
  containers:
    - command:
        - kube-apiserver
        - --anonymous-auth=false
        - --authorization-mode=Node,RBAC
        - --enable-admission-plugins=NodeRestriction
        - --insecure-port=0
        - --profiling=false
        - --audit-log-path=/var/log/kubernetes/audit/audit.log
```

### 1.27.2 Network-level Restrictions

```
# Restrict API server access with NetworkPolicy
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: restrict-api-access
  namespace: kube-system
spec:
  podSelector:
```

```

matchLabels:
  component: kube-apiserver
policyTypes:
- Ingress
ingress:
- from:
  - ipBlock:
    cidr: 10.0.0.0/8

```

## 1.28 Admission Controllers

### 1.28.1 Important Security Admission Controllers

| Controller                 | Description                     |
|----------------------------|---------------------------------|
| NodeRestriction            | Limits kubelet permissions      |
| PodSecurity                | Enforces Pod Security Standards |
| ImagePolicyWebhook         | External image validation       |
| ValidatingAdmissionWebhook | Custom validation               |
| MutatingAdmissionWebhook   | Custom mutation                 |

### 1.28.2 Enable Admission Controllers

```

# /etc/kubernetes/manifests/kube-apiserver.yaml
spec:
  containers:
  - command:
    - kube-apiserver
    - --enable-admission-plugins=NodeRestriction,PodSecurity

```


## 1.29 Key Concepts to Remember

1. **RBAC** - Role, ClusterRole, RoleBinding, ClusterRoleBinding
2. **Least Privilege** - Grant minimum required permissions
3. **ServiceAccount** - Disable auto-mount when not needed
4. **Anonymous Auth** - Disable on API server and kubelet
5. **Admission Controllers** - NodeRestriction, PodSecurity

## 1.30 Practice Questions

1. How do you check if a user can perform an action?
2. What is the difference between Role and ClusterRole?
3. How do you disable ServiceAccount token auto-mount?
4. What admission controller restricts kubelet permissions?
5. How do you upgrade a Kubernetes cluster securely?

## 1.31 System Hardening



[Download PDF Version](#)

This domain covers hardening the underlying host systems and using kernel security features.

## 1.32 AppArmor

### 1.32.1 AppArmor Basics

```
# Check AppArmor status
aa-status
systemctl status apparmor

# List loaded profiles
cat /sys/kernel/security/apparmor/profiles
```

```

# Load a profile
apparmor_parser -q /etc/apparmor.d/my-profile

# Reload a profile
apparmor_parser -r /etc/apparmor.d/my-profile

```

### 1.32.2 AppArmor Profile Example

```

# /etc/apparmor.d/k8s-deny-write
#include <tunables/global>

profile k8s-deny-write flags=(attach_disconnected) {
    #include <abstractions/base>

    file,
    # Deny all file writes
    deny /** w,
}

```

### 1.32.3 Use AppArmor in Pod

```

apiVersion: v1
kind: Pod
metadata:
  name: apparmor-pod
  annotations:
    container.apparmor.security.beta.kubernetes.io/nginx:
      localhost/k8s-deny-write
spec:
  containers:
  - name: nginx
    image: nginx

```

### 1.32.4 AppArmor Profile Modes

| Mode       | Description                |
|------------|----------------------------|
| enforce    | Enforces the policy        |
| complain   | Logs violations but allows |
| unconfined | No restrictions            |

## 1.33 seccomp (Secure Computing Mode)

### 1.33.1 seccomp Profile Location

```

# Default kubelet seccomp profile path
/var/lib/kubelet/seccomp/

```

```
# Create profile directory
mkdir -p /var/lib/kubelet/seccomp/profiles
```

### 1.33.2 seccomp Profile Example

```
{
  "defaultAction": "SCMP_ACT_ERRNO",
  "architectures": [
    "SCMP_ARCH_X86_64"
  ],
  "syscalls": [
    {
      "names": [
        "accept4",
        "bind",
        "clone",
        "close",
        "connect",
        "epoll_create1",
        "epoll_ctl",
        "epoll_pwait",
        "execve",
        "exit_group",
        "fcntl",
        "fstat",
        "futex",
        "getdents64",
        "getpid",
        "getrandom",
        "getsockname",
        "getsockopt",
        "listen",
        "mmap",
        "mprotect",
        "nanosleep",
        "newfstatat",
        "openat",
        "read",
        "recvfrom",
        "rt_sigaction",
        "rt_sigprocmask",
        "sendto",
        "set_tid_address",
        "setsockopt",
        "socket",
        "write"
      ],
      "action": "SCMP_ACT_ALLOW"
    }
  ]
}
```

```
    ]  
}
```

### 1.33.3 Use seccomp in Pod

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: seccomp-pod  
spec:  
  securityContext:  
    seccompProfile:  
      type: Localhost  
      localhostProfile: profiles/audit.json  
  containers:  
  - name: nginx  
    image: nginx
```

### 1.33.4 seccomp Profile Types

| Type           | Description                       |
|----------------|-----------------------------------|
| RuntimeDefault | Container runtime default profile |
| Unconfined     | No seccomp filtering              |
| Localhost      | Custom profile from node          |

### 1.33.5 RuntimeDefault seccomp

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: runtime-default-pod  
spec:  
  securityContext:  
    seccompProfile:  
      type: RuntimeDefault  
  containers:  
  - name: nginx  
    image: nginx
```

## 1.34 Minimize Host OS Footprint

### 1.34.1 Reduce Attack Surface

```
# Remove unnecessary packages  
apt-get remove --purge <package>  
  
# Disable unnecessary services  
systemctl disable <service>
```

```
systemctl stop <service>

# Check listening ports
netstat -tlnp
ss -tlnp

# Check running processes
ps aux
```

### 1.34.2 Secure SSH

```
# /etc/ssh/sshd_config
PermitRootLogin no
PasswordAuthentication no
PubkeyAuthentication yes
PermitEmptyPasswords no
X11Forwarding no
MaxAuthTries 3
```

### 1.34.3 File System Security

```
# Set proper permissions
chmod 600 /etc/kubernetes/admin.conf
chmod 600 /etc/kubernetes/pki/*.key

# Check file permissions
ls -la /etc/kubernetes/
ls -la /etc/kubernetes/pki/
```

## 1.35 Limit Node Access

### 1.35.1 Restrict SSH Access

```
# Allow only specific users
# /etc/ssh/sshd_config
AllowUsers admin

# Use SSH keys only
PasswordAuthentication no
```

### 1.35.2 Use Network Policies for Node Access

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: restrict-node-access
  namespace: kube-system
spec:
  podSelector: {}
```

```
policyTypes:  
  - Ingress  
ingress:  
  - from:  
    - ipBlock:  
      cidr: 10.0.0.0/8
```

## 1.36 Linux Capabilities

### 1.36.1 Drop All Capabilities

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: drop-caps-pod  
spec:  
  containers:  
    - name: nginx  
      image: nginx  
      securityContext:  
        capabilities:  
          drop:  
            - ALL
```

### 1.36.2 Add Specific Capabilities

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: net-admin-pod  
spec:  
  containers:  
    - name: nginx  
      image: nginx  
      securityContext:  
        capabilities:  
          drop:  
            - ALL  
          add:  
            - NET_BIND_SERVICE
```

### 1.36.3 Common Capabilities

| Capability       | Description                       |
|------------------|-----------------------------------|
| NET_BIND_SERVICE | Bind to ports < 1024              |
| NET_ADMIN        | Network administration            |
| SYS_ADMIN        | System administration (dangerous) |
| SYS_PTRACE       | Trace processes                   |

| Capability | Description           |
|------------|-----------------------|
| CHOWN      | Change file ownership |

## 1.37 Kernel Hardening

### 1.37.1 sysctl Settings

```
# /etc/sysctl.d/99-kubernetes-cis.conf

# Disable IP forwarding (if not needed)
net.ipv4.ip_forward = 0

# Disable source routing
net.ipv4.conf.all.accept_source_route = 0
net.ipv4.conf.default.accept_source_route = 0

# Enable SYN cookies
net.ipv4.tcp_syncookies = 1

# Disable ICMP redirects
net.ipv4.conf.all.accept_redirects = 0
net.ipv4.conf.default.accept_redirects = 0

# Apply settings
sysctl --system
```

## 1.38 IAM Roles (Cloud)

### 1.38.1 AWS IAM for Nodes

```
# Restrict node IAM role permissions
# Only allow necessary EC2 and ECR actions
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": [
                "ec2:DescribeInstances",
                "ecr:GetAuthorizationToken",
                "ecr:BatchGetImage"
            ],
            "Resource": "*"
        }
    ]
}
```

## 1.38.2 Pod Identity (IRSA)

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: my-sa
  annotations:
    eks.amazonaws.com/role-arn: arn:aws:iam::123456789:role/my-
      role
```

## 1.39 Key Concepts to Remember

1. **AppArmor** - Linux security module for access control
2. **seccomp** - System call filtering
3. **Capabilities** - Drop ALL, add only needed
4. **Minimize footprint** - Remove unnecessary packages/services
5. **Kernel hardening** - sysctl security settings

## 1.40 Practice Questions


1. How do you apply an AppArmor profile to a container?
2. What is the difference between seccomp RuntimeDefault and Localhost?
3. How do you drop all capabilities from a container?
4. Where are seccomp profiles stored on the node?
5. How do you check AppArmor status on a node?

---

[← Previous: Cluster Hardening](#) | [Back to CKS Overview](#) | [Next: Minimize Microservice Vulnerabilities →](#)

---

## 1.41 Minimize Microservice Vulnerabilities



[Download PDF Version](#)

This domain covers securing containerized applications and implementing Pod security.

## 1.42 Pod Security Standards

### 1.42.1 Pod Security Admission

```
# Apply to namespace
apiVersion: v1
kind: Namespace
metadata:
  name: secure-ns
  labels:
    pod-security.kubernetes.io/enforce: restricted
    pod-security.kubernetes.io/audit: restricted
    pod-security.kubernetes.io/warn: restricted
```

## 1.42.2 Security Levels

| Level      | Description   |
|------------|---|
| privileged | Unrestricted, allows all                                    |
| baseline   | Minimally restrictive, prevents known privilege escalations |
| restricted | Heavily restricted, follows security best practices         |

## 1.42.3 Restricted Pod Example

```
apiVersion: v1
kind: Pod
metadata:
  name: restricted-pod
spec:
  securityContext:
    runAsNonRoot: true
    runAsUser: 1000
    runAsGroup: 3000
    fsGroup: 2000
    seccompProfile:
      type: RuntimeDefault
  containers:
  - name: app
    image: nginx
    securityContext:
      allowPrivilegeEscalation: false
      readOnlyRootFilesystem: true
      capabilities:
        drop:
        - ALL
    volumeMounts:
    - name: tmp
      mountPath: /tmp
    - name: cache
      mountPath: /var/cache/nginx
    - name: run
      mountPath: /var/run
  volumes:
  - name: tmp
    emptyDir: {}
  - name: cache
    emptyDir: {}
  - name: run
    emptyDir: {}
```

## 1.43 Security Context

### 1.43.1 Pod-level Security Context

```
apiVersion: v1
kind: Pod
metadata:
  name: secure-pod
spec:
  securityContext:
    runAsUser: 1000
    runAsGroup: 3000
    fsGroup: 2000
    runAsNonRoot: true
    seccompProfile:
      type: RuntimeDefault
  containers:
    - name: app
      image: nginx
```

### 1.43.2 Container-level Security Context

```
apiVersion: v1
kind: Pod
metadata:
  name: container-secure
spec:
  containers:
    - name: app
      image: nginx
      securityContext:
        runAsUser: 1000
        runAsNonRoot: true
        readOnlyRootFilesystem: true
        allowPrivilegeEscalation: false
      capabilities:
        drop:
          - ALL
        add:
          - NET_BIND_SERVICE
```

### 1.43.3 Security Context Fields

| Field        | Level         | Description          |
|--------------|---------------|----------------------|
| runAsUser    | Pod/Container | UID to run as        |
| runAsGroup   | Pod/Container | GID to run as        |
| runAsNonRoot | Pod/Container | Must run as non-root |
| fsGroup      | Pod           | Group for volumes    |

| Field                    | Level         | Description                  |
|--------------------------|---------------|------------------------------|
| readOnlyRootFilesystem   | Container     | Read-only root FS            |
| allowPrivilegeEscalation | Container     | Prevent privilege escalation |
| capabilities             | Container     | Linux capabilities           |
| seccompProfile           | Pod/Container | seccomp profile              |

## 1.44 OPA Gatekeeper

### 1.44.1 Install Gatekeeper

```
kubectl apply -f https://raw.githubusercontent.com/open-policy-agent/gatekeeper/master/deploy/gatekeeper.yaml
```

### 1.44.2 Constraint Template

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8srequiredlabels
spec:
  crd:
    spec:
      names:
        kind: K8sRequiredLabels
      validation:
        openAPIV3Schema:
          type: object
          properties:
            labels:
              type: array
              items:
                type: string
  targets:
    - target: admission.k8s.gatekeeper.sh
      rego: |
        package k8srequiredlabels

        violation[{"msg": msg}] {
          provided := {label |
            input.review.object.metadata.labels[label]}
          required := {label | label :=
            input.parameters.labels[_]}
          missing := required - provided
          count(missing) > 0
          msg := sprintf("Missing required labels: %v", [missing])
        }
```

### 1.44.3 Constraint

```
apiVersion: constraints.gatekeeper.sh/v1beta1
kind: K8sRequiredLabels
metadata:
  name: require-team-label
spec:
  match:
    kinds:
      - apiGroups: [""]
        kinds: ["Pod"]
  parameters:
    labels: ["team"]
```

### 1.44.4 Deny Privileged Containers

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8sdenyprivileged
spec:
  crd:
    spec:
      names:
        kind: K8sDenyPrivileged
  targets:
    - target: admission.k8s.gatekeeper.sh
      rego: |
        package k8sdenyprivileged

        violation[{"msg": msg}] {
          container := input.review.object.spec.containers[_]
          container.securityContext.privileged == true
          msg := sprintf("Privileged containers are not allowed:
%v", [container.name])
        }
```

## 1.45 Secrets Management

### 1.45.1 Create Secrets

```
# From literal
kubectl create secret generic db-secret \
--from-literal=username=admin \
--from-literal=password=secret123

# From file
kubectl create secret generic tls-secret \  

```

```
--from-file=tls.crt \
--from-file=tls.key
```

### 1.45.2 Use Secrets as Environment Variables

```
apiVersion: v1
kind: Pod
metadata:
  name: secret-env-pod
spec:
  containers:
    - name: app
      image: nginx
      env:
        - name: DB_PASSWORD
          valueFrom:
            secretKeyRef:
              name: db-secret
              key: password
```

### 1.45.3 Use Secrets as Volumes

```
apiVersion: v1
kind: Pod
metadata:
  name: secret-vol-pod
spec:
  containers:
    - name: app
      image: nginx
      volumeMounts:
        - name: secret-volume
          mountPath: /etc/secrets
          readOnly: true
  volumes:
    - name: secret-volume
      secret:
        secretName: db-secret
        defaultMode: 0400
```

### 1.45.4 Encrypt Secrets at Rest

```
# /etc/kubernetes/encryption-config.yaml
apiVersion: apiserver.config.k8s.io/v1
kind: EncryptionConfiguration
resources:
  - resources:
    - secrets
  providers:
    - aescbc:
      keys:
```

```
    - name: key1
      secret: <base64-encoded-32-byte-key>
    - identity: {}
```

## 1.46 Container Runtime Sandboxing

### 1.46.1 gVisor (runsc)

```
apiVersion: node.k8s.io/v1
kind: RuntimeClass
metadata:
  name: gvisor
handler: runsc
---
apiVersion: v1
kind: Pod
metadata:
  name: sandboxed-pod
spec:
  runtimeClassName: gvisor
  containers:
    - name: app
      image: nginx
```

### 1.46.2 Kata Containers

```
apiVersion: node.k8s.io/v1
kind: RuntimeClass
metadata:
  name: kata
handler: kata-qemu
---
apiVersion: v1
kind: Pod
metadata:
  name: kata-pod
spec:
  runtimeClassName: kata
  containers:
    - name: app
      image: nginx
```

## 1.47 mTLS with Service Mesh

### 1.47.1 Istio Strict mTLS

```
apiVersion: security.istio.io/v1beta1
kind: PeerAuthentication
```

```
metadata:  
  name: default  
  namespace: istio-system  
spec:  
  mtls:  
    mode: STRICT
```

## 1.48 Key Concepts to Remember

1. **Pod Security Standards** - privileged, baseline, restricted
2. **Security Context** - runAsNonRoot, readOnlyRootFilesystem, capabilities
3. **OPA Gatekeeper** - Policy enforcement
4. **Secrets** - Encrypt at rest, mount as volumes with restrictive permissions
5. **Runtime Sandboxing** - gVisor, Kata Containers

## 1.49 Practice Questions


1. How do you enforce the restricted Pod Security Standard on a namespace?
2. What security context fields prevent privilege escalation?
3. How do you create an OPA Gatekeeper constraint?
4. What is the recommended way to mount secrets in pods?
5. How do you use a sandboxed runtime for a pod?

---

[← Previous: System Hardening](#) | [Back to CKS Overview](#) | [Next: Supply Chain Security →](#)

---

## 1.50 Supply Chain Security



[Download PDF Version](#)

This domain covers securing the software supply chain including image security, static analysis, and image signing.

## 1.51 Image Vulnerability Scanning

### 1.51.1 Trivy

```
# Scan image
trivy image nginx:1.21

# Scan with severity filter
trivy image --severity HIGH,CRITICAL nginx:1.21

# Scan and fail on vulnerabilities
trivy image --exit-code 1 --severity CRITICAL nginx:1.21

# Scan local image
```

```

trivy image --input image.tar

# Output as JSON
trivy image -f json -o results.json nginx:1.21

# Scan filesystem
trivy fs /path/to/project

# Scan Kubernetes resources
trivy k8s --report summary cluster

```

## 1.51.2 Trivy Output Example

```

nginx:1.21 (debian 11.2)
=====
Total: 125 (UNKNOWN: 0, LOW: 85, MEDIUM: 26, HIGH: 12, CRITICAL: 2)

```

| Library Version | Vulnerability | Severity | Installed        |
|-----------------|---------------|----------|------------------|
| libssl1.1       | CVE-2022-0778 | HIGH     | 1.1.1k-1+deb11u1 |
| openssl         | CVE-2022-0778 | HIGH     | 1.1.1k-1+deb11u1 |

## 1.51.3 Integrate Trivy in CI/CD

```

# GitHub Actions example
- name: Scan image
  uses: aquasecurity/trivy-action@master
  with:
    image-ref: 'myimage:${{ github.sha }}'
    format: 'table'
    exit-code: '1'
    severity: 'CRITICAL,HIGH'

```

## 1.52 Image Signing and Verification

### 1.52.1 Cosign

```

# Generate key pair
cosign generate-key-pair

# Sign image
cosign sign --key cosign.key myregistry/myimage:v1

# Verify image

```

```

cosign verify --key cosign.pub myregistry/myimage:v1

# Sign with keyless (OIDC)
cosign sign myregistry/myimage:v1

# Verify keyless signature
cosign verify myregistry/myimage:v1

```

## 1.52.2 Image Policy Webhook

```

apiVersion: admissionregistration.k8s.io/v1
kind: ValidatingWebhookConfiguration
metadata:
  name: image-policy-webhook
webhooks:
- name: image-policy.example.com
  rules:
  - apiGroups: [""]
    apiVersions: ["v1"]
    operations: ["CREATE", "UPDATE"]
    resources: ["pods"]
  clientConfig:
    service:
      name: image-policy-webhook
      namespace: kube-system
      path: "/validate"
      caBundle: <base64-encoded-ca-cert>
  admissionReviewVersions: ["v1"]
  sideEffects: None

```

## 1.53 Minimize Base Image Footprint

### 1.53.1 Use Minimal Base Images

```

# Bad: Full OS image
FROM ubuntu:22.04

# Better: Slim image
FROM python:3.11-slim

# Best: Distroless
FROM gcr.io/distroless/python3

# Best: Scratch (for static binaries)
FROM scratch

```

### 1.53.2 Multi-stage Builds

```

# Build stage
FROM golang:1.21 AS builder

```

```

WORKDIR /app
COPY .
RUN CGO_ENABLED=0 GOOS=linux go build -o myapp

# Runtime stage
FROM gcr.io/distroless/static:nonroot
COPY --from=builder /app/myapp /myapp
USER nonroot:nonroot
ENTRYPOINT ["/myapp"]

```

### 1.53.3 Image Best Practices

| Practice            | Description                    |
|---------------------|--------------------------------|
| Use specific tags   | Avoid latest, use version tags |
| Use minimal base    | distroless, alpine, scratch    |
| Multi-stage builds  | Reduce final image size        |
| Run as non-root     | Use USER instruction           |
| No secrets in image | Use runtime secrets            |
| Scan regularly      | Integrate scanning in CI/CD    |

## 1.54 Secure Image Registries

### 1.54.1 Private Registry Authentication

```

# Create registry secret
kubectl create secret docker-registry regcred \
--docker-server=myregistry.io \
--docker-username=user \
--docker-password=pass \
--docker-email=user@example.com

```

### 1.54.2 Use ImagePullSecrets

```

apiVersion: v1
kind: Pod
metadata:
  name: private-image-pod
spec:
  containers:
  - name: app
    image: myregistry.io/myimage:v1
  imagePullSecrets:
  - name: regcred

```

### 1.54.3 ServiceAccount with ImagePullSecrets

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: my-sa
imagePullSecrets:
- name: regcred
```

## 1.55 Allowlist Registries

### 1.55.1 OPA Gatekeeper Policy

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8sallowedrepos
spec:
  crd:
    spec:
      names:
        kind: K8sAllowedRepos
      validation:
        openAPIV3Schema:
          type: object
          properties:
            repos:
              type: array
              items:
                type: string
      targets:
        - target: admission.k8s.gatekeeper.sh
          rego: |
            package k8sallowedrepos

            violation[{"msg": msg}] {
              container := input.review.object.spec.containers[_]
              not startswith(container.image,
              input.parameters.repos[_])
              msg := sprintf("Container image %v is not from an
              allowed registry", [container.image])
            }

            ---
apiVersion: constraints.gatekeeper.sh/v1beta1
kind: K8sAllowedRepos
metadata:
  name: allowed-repos
spec:
  match:
    kinds:
      - apiGroups: ["""]
```

```

        kinds: ["Pod"]
parameters:
repos:
- "gcr.io/my-project/"
- "myregistry.io/"

```

## 1.56 Static Analysis

### 1.56.1 Kubesec

```

# Scan Kubernetes manifest
kubesec scan pod.yaml

# Scan with HTTP API
curl -sSX POST --data-binary @pod.yaml https://v2.kubesec.io/scan

```

### 1.56.2 Kubesec Output

```
{
  "score": 0,
  "scoring": {
    "critical": [
      {
        "id": "CapSysAdmin",
        "selector": "containers[] .securityContext .capabilities .add == SYS_ADMIN",
        "reason": "CAP_SYS_ADMIN is the most privileged capability"
      }
    ],
    "advise": [
      {
        "id": "ApparmorAny",
        "selector": ".metadata .annotations .\"container.apparmor.security.beta.kubernetes.io/nginx\"",
        "reason": "Well defined AppArmor policies may provide greater protection"
      }
    ]
  }
}
```

### 1.56.3 Conftest (OPA for Config Files)

```

# Test Kubernetes manifests
conftest test deployment.yaml

```

```
# With custom policy
conftest test --policy ./policy deployment.yaml
```

## 1.56.4 Example Conftest Policy

```
# policy/deployment.rego
package main

deny[msg] {
    input.kind == "Deployment"
    not input.spec.template.spec.securityContext.runAsNonRoot
    msg := "Containers must run as non-root"
}

deny[msg] {
    input.kind == "Deployment"
    container := input.spec.template.spec.containers[_]
    not container.resources.limits
    msg := sprintf("Container %v must have resource limits",
    [container.name])
}
```

## 1.57 Dockerfile Security

### 1.57.1 Dockerfile Best Practices

```
# Use specific version
FROM python:3.11-slim-bookworm

# Create non-root user
RUN groupadd -r appgroup && useradd -r -g appgroup appuser

# Set working directory
WORKDIR /app

# Copy only necessary files
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt

COPY --chown=appuser:appgroup .

# Switch to non-root user
USER appuser

# Use ENTRYPPOINT for main command
ENTRYPOINT ["python", "app.py"]
```

## 1.57.2 Hadolint (Dockerfile Linter)

```
# Lint Dockerfile
hadolint Dockerfile

# Ignore specific rules
hadolint --ignore DL3008 Dockerfile
```

## 1.58 Key Concepts to Remember

1. **Trivy** - Image vulnerability scanning
2. **Cosign** - Image signing and verification
3. **Minimal images** - distroless, alpine, scratch
4. **Multi-stage builds** - Reduce image size
5. **Registry allowlisting** - OPA Gatekeeper policies

## 1.59 Practice Questions


1. How do you scan an image for HIGH and CRITICAL vulnerabilities with Trivy?
2. How do you sign an image with Cosign?
3. What is the benefit of using distroless base images?
4. How do you create an OPA policy to allowlist registries?
5. What tool lints Dockerfiles for security issues?

---

[← Previous: Minimize Microservice Vulnerabilities](#) | [Back to CKS Overview](#) | [Next: Monitoring, Logging and Runtime Security →](#)

---

## 1.60 Monitoring, Logging and Runtime Security



[Download PDF Version](#)

This domain covers runtime security monitoring, audit logging, behavioral analysis, and threat detection in Kubernetes environments.

## 1.61 Overview

Runtime security focuses on detecting and preventing threats while workloads are running. Key areas include:

- Behavioral analysis and anomaly detection
- Audit logging for compliance and forensics
- Container immutability enforcement
- System call monitoring

## 1.62 Falco - Runtime Threat Detection


Falco is a CNCF project that provides runtime security by monitoring system calls and detecting anomalous behavior.

## 1.62.1 Install Falco

```
# Using Helm
helm repo add falcosecurity https://falcosecurity.github.io/charts
helm repo update
helm install falco falcosecurity/falco \
    --namespace falco \
    --create-namespace \
    --set falcosidekick.enabled=true

# Verify installation
kubectl get pods -n falco
kubectl logs -n falco -l app.kubernetes.io/name=falco
```

## 1.62.2 Falco Architecture



### 1.62.3 Falco Rules Structure

```
tags: [<tag1>, <tag2>]      # Optional categorization
enabled: true/false          # Enable/disable rule
```

#### 1.62.4 Common Falco Rules

```
# Detect shell spawning in containers
- rule: Terminal shell in container
  desc: A shell was used as the entrypoint/exec point into a
        container
  condition: >
    spawned_process and container
    and shell_procs
    and proc.tty != 0
    and container_entrypoint
  output: >
    Shell spawned in container
    (user=%user.name user_loginuid=%user.loginuid %container.info
     shell=%proc.name parent=%proc.pname cmdline=%proc.cmdline
     terminal=%proc.tty container_id=%container.id
     image=%container.image.repository)
  priority: WARNING
  tags: [container, shell, mitre_execution]

# Detect writes to sensitive directories
- rule: Write below etc
  desc: An attempt to write to /etc directory
  condition: >
    write and container
    and fd.directory = /etc
  output: >
    File written to /etc in container
    (user=%user.name command=%proc.cmdline
     file=%fd.name container=%container.name
     image=%container.image.repository)
  priority: ERROR
  tags: [filesystem, mitre_persistence]

# Detect package management in containers
- rule: Package management in container
  desc: Package management tool executed in container
  condition: >
    spawned_process and container
    and package_mgmt_procs
  output: >
    Package management process launched in container
    (user=%user.name command=%proc.cmdline
     container=%container.name)
  priority: ERROR
  tags: [process, software_mgmt]

# Detect sensitive file access
- rule: Read sensitive file
  desc: Sensitive file was read
```

```

condition: >
  open_read and container
  and (fd.name startswith /etc/shadow or
       fd.name startswith /etc/passwd or
       fd.name startswith /etc/sudoers)
output: >
  Sensitive file opened for reading
  (user=%user.name file=%fd.name container=%container.name)
priority: WARNING
tags: [filesystem, mitre_credential_access]

# Detect outbound connections to unusual ports
- rule: Unexpected outbound connection
  desc: Outbound connection to non-standard port
  condition: >
    outbound and container
    and not (fd.sport in (80, 443, 53, 8080, 8443))
  output: >
    Unexpected outbound connection
    (command=%proc.cmdline connection=%fd.name
     container=%container.name)
  priority: NOTICE
  tags: [network, mitre_exfiltration]

```

### 1.62.5 Custom Falco Rules

```

# /etc/falco/rules.d/custom-rules.yaml
- rule: Crypto mining detected
  desc: Detect crypto mining processes
  condition: >
    spawned_process and container
    and (proc.name in (xmrig, minerd, cpuminer, cgminer))
  output: >
    Crypto mining process detected
    (user=%user.name command=%proc.cmdline
     container=%container.name)
  priority: CRITICAL
  tags: [cryptomining, mitre_resource_hijacking]

- rule: Kubectl exec detected
  desc: Detect kubectl exec into containers
  condition: >
    spawned_process and container
    and proc.name = "runc"
    and proc.cmdline contains "exec"
  output: >
    Kubectl exec detected (container=%container.name
      cmdline=%proc.cmdline)
  priority: WARNING
  tags: [exec, mitre_execution]

```

## 1.62.6 Falco Commands

```
# Check Falco status
systemctl status falco

# View Falco logs
kubectl logs -n falco -l app.kubernetes.io/name=falco -f

# Test Falco rules
kubectl exec -it <pod> -- /bin/sh # Should trigger shell alert

# Reload rules
kill -1 $(pidof falco)

# Check rule syntax
falco -V /etc/falco/rules.d/custom-rules.yaml
```

## 1.63 Kubernetes Audit Logging

Audit logging records all requests to the Kubernetes API server for security analysis and compliance.

### 1.63.1 Audit Stages

| Stage                   | Description                              |
|-------------------------|--|
| <b>RequestReceived</b>  | Event generated when request is received |
| <b>ResponseStarted</b>  | Response headers sent, body not yet sent |
| <b>ResponseComplete</b> | Response body completed                  |
| <b>Panic</b>            | Events generated when panic occurs       |

### 1.63.2 Audit Levels

| Level                  | Description                                |
|------------------------|--|
| <b>None</b>            | Don't log events                           |
| <b>Metadata</b>        | Log request metadata only                  |
| <b>Request</b>         | Log metadata and request body              |
| <b>RequestResponse</b> | Log metadata, request, and response bodies |

### 1.63.3 Comprehensive Audit Policy

```
apiVersion: audit.k8s.io/v1
kind: Policy
rules:
  # Don't log read-only endpoints
  - level: None
    nonResourceURLs:
```

```
- /healthz*
- /version
- /swagger*
- /readyz*
- /livez*

# Don't log watch requests (too verbose)
- level: None
  verbs: ["watch"]
  resources:
    - group: ""
      resources: ["events"]

# Log secret access at Metadata level (don't log secret data)
- level: Metadata
  resources:
    - group: ""
      resources: ["secrets", "configmaps"]

# Log pod exec/attach at RequestResponse level
- level: RequestResponse
  resources:
    - group: ""
      resources: ["pods/exec", "pods/attach", "pods/portforward"]

# Log authentication events
- level: RequestResponse
  resources:
    - group: "authentication.k8s.io"
      resources: ["tokenreviews"]

# Log authorization events
- level: RequestResponse
  resources:
    - group: "authorization.k8s.io"
      resources: ["subjectaccessreviews",
                  "selfsubjectaccessreviews"]

# Log RBAC changes
- level: RequestResponse
  resources:
    - group: "rbac.authorization.k8s.io"
      resources: ["roles", "rolebindings", "clusterroles",
                  "clusterrolebindings"]

# Log node and namespace changes
- level: RequestResponse
  resources:
    - group: ""
      resources: ["nodes", "namespaces"]

# Log service account token creation
```

```

- level: RequestResponse
  resources:
    - group: ""
      resources: ["serviceaccounts/token"]

# Default: log at Metadata level
- level: Metadata
  omitStages:
    - RequestReceived

```

#### 1.63.4 Enable Audit Logging on API Server

```

# /etc/kubernetes/manifests/kube-apiserver.yaml
apiVersion: v1
kind: Pod
metadata:
  name: kube-apiserver
  namespace: kube-system
spec:
  containers:
    - command:
        - kube-apiserver
        # Audit configuration
        - --audit-policy-file=/etc/kubernetes/audit/audit-policy.yaml
        - --audit-log-path=/var/log/kubernetes/audit/audit.log
        - --audit-log-maxage=30
        - --audit-log-maxbackup=10
        - --audit-log-maxsize=100
        # Optional: webhook backend
        # - --audit-webhook-config-file=/etc/kubernetes/audit/webhook-
          config.yaml
  volumeMounts:
    - mountPath: /etc/kubernetes/audit
      name: audit-config
      readOnly: true
    - mountPath: /var/log/kubernetes/audit
      name: audit-log
  volumes:
    - hostPath:
        path: /etc/kubernetes/audit
        type: DirectoryOrCreate
        name: audit-config
    - hostPath:
        path: /var/log/kubernetes/audit
        type: DirectoryOrCreate
        name: audit-log

```

#### 1.63.5 Analyze Audit Logs

```

# View all audit events
cat /var/log/kubernetes/audit/audit.log | jq .

```

```

# Find delete operations
cat /var/log/kubernetes/audit/audit.log | jq
    'select(.verb=="delete")'

# Find secret access
cat /var/log/kubernetes/audit/audit.log | jq
    'select(.objectRef.resource=="secrets")'

# Find failed requests
cat /var/log/kubernetes/audit/audit.log | jq
    'select(.responseStatus.code >= 400)'

# Find requests by user
cat /var/log/kubernetes/audit/audit.log | jq
    'select(.user.username=="system:admin")'

# Find pod exec events
cat /var/log/kubernetes/audit/audit.log | jq
    'select(.objectRef.subresource=="exec")'

# Find RBAC changes
cat /var/log/kubernetes/audit/audit.log | jq
    'select(.objectRef.apiGroup=="rbac.authorization.k8s.io")'

# Count events by verb
cat /var/log/kubernetes/audit/audit.log | jq -r '.verb' | sort |
    uniq -c | sort -rn

```

### 1.63.6 Audit Log Entry Structure

```
{
    "kind": "Event",
    "apiVersion": "audit.k8s.io/v1",
    "level": "RequestResponse",
    "auditID": "unique-id",
    "stage": "ResponseComplete",
    "requestURI": "/api/v1/namespaces/default/pods",
    "verb": "create",
    "user": {
        "username": "admin",
        "groups": ["system:masters"]
    },
    "sourceIPs": ["192.168.1.100"],
    "objectRef": {
        "resource": "pods",
        "namespace": "default",
        "name": "nginx",
        "apiVersion": "v1"
    },
    "responseStatus": {
        "code": 201
    },
    "requestObject": { }
}
```

```
"responseObject": { },
"requestReceivedTimestamp": "2024-01-15T10:30:00.000000Z",
"stageTimestamp": "2024-01-15T10:30:00.100000Z"
}
```

## 1.64 Container Immutability

Immutable containers prevent runtime modifications, reducing attack surface.

### 1.64.1 Read-Only Root Filesystem

```
apiVersion: v1
kind: Pod
metadata:
  name: immutable-pod
spec:
  containers:
    - name: app
      image: nginx:1.21
      securityContext:
        readOnlyRootFilesystem: true
        allowPrivilegeEscalation: false
        runAsNonRoot: true
        runAsUser: 1000
      capabilities:
        drop:
          - ALL
  volumeMounts:
    # Mount writable directories as needed
    - name: tmp
      mountPath: /tmp
    - name: var-run
      mountPath: /var/run
    - name: var-cache-nginx
      mountPath: /var/cache/nginx
  volumes:
    - name: tmp
      emptyDir: {}
    - name: var-run
      emptyDir: {}
    - name: var-cache-nginx
      emptyDir: {}
```

### 1.64.2 Enforce Immutability with Pod Security Standards

```
apiVersion: v1
kind: Namespace
metadata:
  name: production
  labels:
```

```
pod-security.kubernetes.io/enforce: restricted  
pod-security.kubernetes.io/audit: restricted  
pod-security.kubernetes.io/warn: restricted
```

## 1.65 Sysdig and Other Tools

### 1.65.1 Sysdig for Troubleshooting

```
# Capture system calls  
sysdig -c topprocs_cpu  
  
# Monitor container activity  
sysdig -c topcontainers_cpu  
  
# Watch file opens in container  
sysdig -c echo_fds container.name=nginx  
  
# Monitor network connections  
sysdig -c topconns  
  
# Capture to file for analysis  
sysdig -w capture.scap  
sysdig -r capture.scap
```


## 1.66 Key Concepts Summary

| Concept                       | Description                            | Tool/Feature           |
|-------------------------------|--|------------------------|
| <b>Runtime Detection</b>      | Monitor running containers for threats | Falco                  |
| <b>Audit Logging</b>          | Record API server activity             | Kubernetes Audit       |
| <b>Immutability</b>           | Prevent runtime modifications          | readOnlyRootFilesystem |
| <b>Behavioral Analysis</b>    | Detect anomalous behavior              | Falco rules            |
| <b>System Call Monitoring</b> | Track kernel-level activity            | eBPF, Sysdig           |

## 1.67 Practice Exercises

1. Install Falco and trigger a shell alert
2. Create a custom Falco rule to detect file writes to /tmp
3. Configure audit logging to capture secret access
4. Analyze audit logs to find failed authentication attempts
5. Create an immutable pod with read-only filesystem

## 1.68 Sample Practice Questions



[Download PDF Version](#)

**Disclaimer:** These are sample practice questions created for study purposes only. They are NOT actual exam questions and are designed to help you test your understanding of CKS concepts.

## 1.69 Practice Resources

- [Killercoda CKS Scenarios](#) ★ Free hands-on practice
  - [killer.sh CKS Simulator](#) - Included with exam registration
-

## 1.70 Section 1: Cluster Setup (10%)

### 1.70.1 Question 1.1 - Network Policy

Create a NetworkPolicy named deny-all in namespace secure that denies all ingress and egress traffic.

Show Solution

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: deny-all
  namespace: secure
spec:
  podSelector: {}
  policyTypes:
    - Ingress
    - Egress
```

### 1.70.2 Question 1.2 - CIS Benchmark

Run kube-bench to check the master node against CIS benchmarks.

Show Solution

```
kube-bench run --targets=master
```

---

## 1.71 Section 2: Cluster Hardening (15%)

### 1.71.1 Question 2.1 - RBAC

Create a Role named pod-reader that allows get, list, watch on pods. Bind it to user jane.

Show Solution

```
kubectl create role pod-reader --verb=get,list,watch --
      resource=pods
kubectl create rolebinding read-pods --role=pod-reader --user=jane
```

### 1.71.2 Question 2.2 - ServiceAccount

Create a Pod that doesn't auto-mount the ServiceAccount token.

Show Solution

```
apiVersion: v1
kind: Pod
```

```
metadata:
  name: no-token-pod
spec:
  automountServiceAccountToken: false
  containers:
  - name: nginx
    image: nginx
```

---

## 1.72 Section 3: System Hardening (15%)

### 1.72.1 Question 3.1 - AppArmor

Create a Pod with AppArmor profile k8s-deny-write applied.

Show Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: apparmor-pod
  annotations:
    container.apparmor.security.beta.kubernetes.io/nginx:
      localhost/k8s-deny-write
spec:
  containers:
  - name: nginx
    image: nginx
```

### 1.72.2 Question 3.2 - seccomp

Create a Pod using RuntimeDefault seccomp profile.

Show Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: seccomp-pod
spec:
  securityContext:
    seccompProfile:
      type: RuntimeDefault
  containers:
  - name: nginx
    image: nginx
```

---

## 1.73 Section 4: Minimize Microservice Vulnerabilities (20%)

### 1.73.1 Question 4.1 - Security Context

Create a Pod that runs as non-root with read-only filesystem and drops all capabilities.

Show Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: secure-pod
spec:
  securityContext:
    runAsNonRoot: true
    runAsUser: 1000
  containers:
  - name: nginx
    image: nginx
    securityContext:
      readOnlyRootFilesystem: true
      allowPrivilegeEscalation: false
    capabilities:
      drop:
      - ALL
```

### 1.73.2 Question 4.2 - Pod Security Standards

Apply the restricted Pod Security Standard to namespace production.

Show Solution

```
kubectl label namespace production \
  pod-security.kubernetes.io/enforce=restricted \
  pod-security.kubernetes.io/audit=restricted \
  pod-security.kubernetes.io/warn=restricted
```

---

## 1.74 Section 5: Supply Chain Security (20%)

### 1.74.1 Question 5.1 - Image Scanning

Scan image nginx:1.21 for HIGH and CRITICAL vulnerabilities using Trivy.

Show Solution

```
trivy image --severity HIGH,CRITICAL nginx:1.21
```

## 1.74.2 Question 5.2 - Allowed Registries

Create an OPA Gatekeeper constraint to only allow images from gcr.io/myproject/.

Show Solution

```
apiVersion: constraints.gatekeeper.sh/v1beta1
kind: K8sAllowedRepos
metadata:
  name: allowed-repos
spec:
  match:
    kinds:
      - apiGroups: [""]
        kinds: ["Pod"]
  parameters:
    repos:
      - "gcr.io/myproject/"
```

---

## 1.75 Section 6: Monitoring, Logging and Runtime Security (20%)

### 1.75.1 Question 6.1 - Audit Logging

Configure API server to log all secret access at Metadata level.

Show Solution

```
apiVersion: audit.k8s.io/v1
kind: Policy
rules:
  - level: Metadata
    resources:
      - group: ""
        resources: ["secrets"]
```

Add to kube-apiserver:

```
--audit-policy-file=/etc/kubernetes/audit-policy.yaml
--audit-log-path=/var/log/kubernetes/audit/audit.log
```

### 1.75.2 Question 6.2 - Falco

Write a Falco rule to detect shell spawning in containers.

Show Solution

```
- rule: Shell in container
  desc: Detect shell spawned in container
```

```
condition: spawned_process and container and shell_procs
output: Shell spawned (user=%user.name
        container=%container.name)
priority: WARNING
```

### 1.75.3 Question 6.3 - Container Immutability

Create a Pod with a read-only root filesystem that can still write to /tmp.

Show Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: immutable-pod
spec:
  containers:
  - name: app
    image: nginx
    securityContext:
      readOnlyRootFilesystem: true
    volumeMounts:
    - name: tmp
      mountPath: /tmp
  volumes:
  - name: tmp
    emptyDir: {}
```

### 1.75.4 Question 6.4 - Analyze Audit Logs

Find all secret access events in the audit log.

Show Solution

```
cat /var/log/kubernetes/audit/audit.log | jq
  'select(.objectRef.resource=="secrets")'

# Or find specific operations on secrets
cat /var/log/kubernetes/audit/audit.log | jq
  'select(.objectRef.resource=="secrets" and .verb=="get")'
```

---

## 1.76 Section 7: Additional Practice Questions

### 1.76.1 Question 7.1 - Ingress TLS

Create an Ingress with TLS termination using a secret named `tls-secret`.

Show Solution

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: secure-ingress
spec:
  tls:
  - hosts:
    - secure.example.com
    secretName: tls-secret
  rules:
  - host: secure.example.com
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: my-service
            port:
              number: 80
```

## 1.76.2 Question 7.2 - RuntimeClass with gVisor

Create a RuntimeClass for gVisor and a Pod that uses it.

Show Solution

```
apiVersion: node.k8s.io/v1
kind: RuntimeClass
metadata:
  name: gvisor
handler: runsc
---
apiVersion: v1
kind: Pod
metadata:
  name: sandboxed-pod
spec:
  runtimeClassName: gvisor
  containers:
  - name: app
    image: nginx
```

## 1.76.3 Question 7.3 - Encrypt etcd Data

Configure encryption at rest for secrets in etcd.

Show Solution

```
# /etc/kubernetes/enc/enc.yaml
apiVersion: apiserver.config.k8s.io/v1
kind: EncryptionConfiguration
```

```

resources:
  - resources:
    - secrets
  providers:
    - aescbc:
      keys:
        - name: key1
          secret: <base64-encoded-32-byte-key>
  - identity: {}

```

Add to kube-apiserver:

```
--encryption-provider-config=/etc/kubernetes/enc/enc.yaml
```

#### 1.76.4 Question 7.4 - Verify Image Signature

Use cosign to verify an image signature.

Show Solution

```

# Verify image signature
cosign verify --key cosign.pub gcr.io/myproject/myimage:v1

# Generate a key pair
cosign generate-key-pair

# Sign an image
cosign sign --key cosign.key gcr.io/myproject/myimage:v1

```

#### 1.76.5 Question 7.5 - Restrict Syscalls with Seccomp

Create a Pod that blocks the chmod syscall using a custom seccomp profile.

Show Solution

```

// /var/lib/kubelet/seccomp/profiles/block-chmod.json
{
  "defaultAction": "SCMP_ACT_ALLOW",
  "syscalls": [
    {
      "names": ["chmod", "fchmod", "fchmodat"],
      "action": "SCMP_ACT_ERRNO"
    }
  ]
}

apiVersion: v1
kind: Pod
metadata:
  name: seccomp-pod
spec:
  securityContext:
    seccompProfile:

```

```
    type: Localhost
    localhostProfile: profiles/block-chmod.json
  containers:
  - name: app
    image: nginx
```

## 1.76.6 Question 7.6 - Network Policy for Database

Create a NetworkPolicy that only allows pods with label `app=backend` to access pods with label `app=database` on port 5432.

Show Solution

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: db-policy
spec:
  podSelector:
    matchLabels:
      app: database
  policyTypes:
  - Ingress
  ingress:
  - from:
    - podSelector:
        matchLabels:
          app: backend
  ports:
  - protocol: TCP
    port: 5432
```

## 1.76.7 Question 7.7 - Admission Controller Webhook

Which admission controllers should be enabled for security?

Show Solution

Essential security admission controllers:

- `NodeRestriction` - Limits kubelet permissions
- `PodSecurity` - Enforces Pod Security Standards
- `AlwaysPullImages` - Forces image pull on every pod start
- `DenyServiceExternalIPs` - Prevents external IP assignment

Enable in kube-apiserver:

```
--enable-admission-
plugins=NodeRestriction,PodSecurity,AlwaysPullImages
```

## 1.76.8 Question 7.8 - Investigate Compromised Pod

A pod is suspected of being compromised. What steps would you take?

Show Solution

```
# 1. Isolate the pod with NetworkPolicy
kubectl apply -f - <<EOF
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: isolate-pod
spec:
  podSelector:
    matchLabels:
      app: compromised
  policyTypes:
  - Ingress
  - Egress
EOF

# 2. Check pod events and logs
kubectl describe pod <pod-name>
kubectl logs <pod-name>

# 3. Check running processes
kubectl exec <pod-name> -- ps aux

# 4. Check network connections
kubectl exec <pod-name> -- netstat -tulpn

# 5. Check audit logs
cat /var/log/kubernetes/audit/audit.log | jq
  'select(.objectRef.name=="<pod-name>")'

# 6. Check Falco alerts
kubectl logs -n falco -l app.kubernetes.io/name=falco | grep <pod-
name>
```

---

## 1.77 Exam Tips

1. **Practice on [KillerCoda](#)** - Free hands-on scenarios
2. **Use [killer.sh](#)** - Included with exam registration
3. **Know Trivy for image scanning** - trivy image --severity HIGH,CRITICAL
4. **Master Network Policies** - Default deny patterns, ingress/egress rules
5. **Understand RBAC deeply** - Roles, ClusterRoles, bindings
6. **Practice with Falco rules** - Rule syntax, common detections
7. **Know audit logging configuration** - Policy levels, API server flags
8. **Understand Pod Security Standards** - privileged, baseline, restricted
9. **Practice seccomp and AppArmor** - Profile application
10. **Know encryption at rest** - etcd encryption configuration

## 1.78 Quick Reference Commands

```
# Image scanning
trivy image nginx:latest

# Check CIS benchmarks
kube-bench run --targets=master

# Create secret
kubectl create secret generic my-secret --from-literal=password=secret

# Apply NetworkPolicy
kubectl apply -f networkpolicy.yaml

# Check RBAC
kubectl auth can-i create pods --as=jane

# View audit logs
cat /var/log/kubernetes/audit/audit.log | jq .

# Check Falco logs
kubectl logs -n falco -l app.kubernetes.io/name=falco
```

---

[← Back to CKS Overview](#)

---