

## 1 CNPE

- 1.1 Table of Contents
- 1.2 Overview
- 1.3 Exam Overview
- 1.4 Exam Domains & Weights
- 1.5 Key Topics
  - 1.5.1 Platform Design
  - 1.5.2 Platform Implementation
  - 1.5.3 Platform Operations
- 1.6 Study Resources
- 1.7 Navigation
- 1.8 Platform Engineering
- 1.9 Overview
- 1.10 Key Concepts
  - 1.10.1 Internal Developer Platform (IDP)
  - 1.10.2 Platform Team Responsibilities
- 1.11 Crossplane
  - 1.11.1 Installation
  - 1.11.2 Composite Resources
  - 1.11.3 Composition
- 1.12 Backstage
  - 1.12.1 Features
  - 1.12.2 Catalog Entity
  - 1.12.3 Software Template
- 1.13 Platform Patterns
  - 1.13.1 Golden Paths
  - 1.13.2 Self-Service Infrastructure
  - 1.13.3 Platform APIs
- 1.14 Best Practices
- 1.15 Sample Practice Questions
- 1.16 Practice Resources
- 1.17 Platform Design (20%)
  - 1.17.1 Question 1
  - 1.17.2 Question 2
- 1.18 Platform Implementation (30%)
  - 1.18.1 Question 3
  - 1.18.2 Question 4
  - 1.18.3 Question 5
- 1.19 Platform Operations (25%)
  - 1.19.1 Question 6
  - 1.19.2 Question 7
- 1.20 Security and Governance (15%)
  - 1.20.1 Question 8
  - 1.20.2 Question 9
- 1.21 Developer Experience (10%)
  - 1.21.1 Question 10
- 1.22 Exam Tips

# 1 CNPE

Generated on: 2026-01-13 15:04:44 Version: 1.0

---

## 1.1 Table of Contents

1. [Overview](#)
  2. [Platform Engineering](#)
  3. [Sample Practice Questions](#)
- 

## 1.2 Overview



The **Certified Cloud Native Platform Engineer (CNPE)** exam demonstrates advanced knowledge of platform engineering, including designing, building, and operating cloud native platforms.

**Note:** This certification is required for Golden Kubestronaut status after March 1st, 2025.

## 1.3 Exam Overview

Detail	Information
<b>Exam Format</b>	Performance-based (hands-on)
<b>Duration</b>	2 hours
<b>Passing Score</b>	66%
<b>Certification Validity</b>	3 years
<b>Cost</b>	\$395 USD
<b>Retake Policy</b>	1 free retake

## 1.4 Exam Domains & Weights

Domain	Weight
Platform Design	20%
Platform Implementation	30%
Platform Operations	25%
Security and Governance	15%
Developer Experience	10%

## 1.5 Key Topics

### 1.5.1 Platform Design

- Architecture patterns
- Multi-tenancy strategies
- API design
- Scalability considerations

### 1.5.2 Platform Implementation

- Kubernetes operators
- Custom controllers
- GitOps implementation
- Infrastructure as Code

### 1.5.3 Platform Operations

- Day-2 operations
- Upgrades and migrations
- Disaster recovery
- Capacity planning

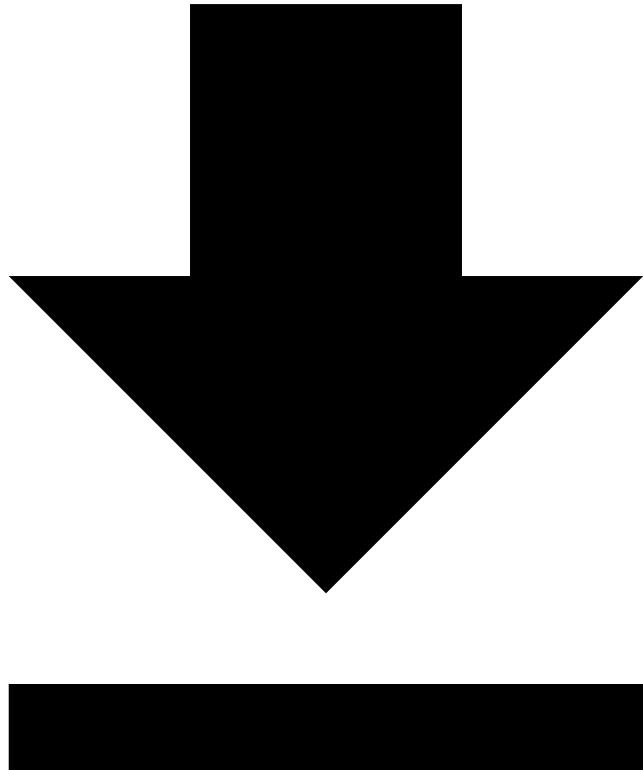
## 1.6 Study Resources

- [Platform Engineering](#)
- [CNCF Platforms White Paper](#)
- [CNPE Curriculum](#)
- [Kubernetes Operators](#)

## 1.7 Navigation

- [Next: Sample Questions →](#)
-

## 1.8 Platform Engineering



[Download PDF Version](#)

Comprehensive guide to platform engineering for CNPE certification.

---

## 1.9 Overview

Platform Engineering focuses on building Internal Developer Platforms (IDPs) that:

- **Abstract complexity** - Hide infrastructure details
  - **Enable self-service** - Developers provision resources independently
  - **Standardize** - Consistent patterns across teams
  - **Automate** - Reduce manual operations
-

## 1.10 Key Concepts

### 1.10.1 Internal Developer Platform (IDP)

An IDP provides:

- Self-service infrastructure provisioning
- Standardized deployment pipelines
- Observability and monitoring
- Security and compliance guardrails

### 1.10.2 Platform Team Responsibilities

- Build and maintain platform components
  - Define golden paths for common workflows
  - Provide documentation and support
  - Measure and improve developer experience
- 

## 1.11 Crossplane

Crossplane extends Kubernetes to manage cloud resources:

### 1.11.1 Installation

```
# Install Crossplane
kubectl create namespace crossplane-system
helm repo add crossplane-stable https://charts.crossplane.io/
    stable
helm install crossplane crossplane-stable/crossplane -n
    crossplane-system

# Install AWS provider
kubectl apply -f - <<EOF
apiVersion: pkg.crossplane.io/v1
kind: Provider
metadata:
  name: provider-aws
spec:
  package: xpkg.upbound.io/crossplane-contrib/provider-aws:v0.40.0
EOF
```

### 1.11.2 Composite Resources

```
apiVersion: apiextensions.crossplane.io/v1
kind: CompositeResourceDefinition
metadata:
  name: databases.platform.example.com
spec:
```

```

group: platform.example.com
names:
  kind: Database
  plural: databases
versions:
- name: v1
  served: true
  referenceable: true
  schema:
    openAPIV3Schema:
      type: object
      properties:
        spec:
          type: object
          properties:
            size:
              type: string
              enum: [small, medium, large]
            engine:
              type: string
              enum: [postgres, mysql]

```

### 1.11.3 Composition

```

apiVersion: apiextensions.crossplane.io/v1
kind: Composition
metadata:
  name: database-aws
spec:
  compositeTypeRef:
    apiVersion: platform.example.com/v1
    kind: Database
  resources:
    - name: rds-instance
      base:
        apiVersion: rds.aws.crossplane.io/v1beta1
        kind: Instance
        spec:
          forProvider:
            region: us-east-1
            dbInstanceClass: db.t3.micro
            engine: postgres
            masterUsername: admin
      patches:
        - fromFieldPath: spec.size
          toFieldPath: spec.forProvider.dbInstanceClass
          transforms:
            - type: map
              map:
                small: db.t3.micro

```

```
medium: db.t3.small
large: db.t3.medium
```

---

## 1.12 Backstage

Backstage is a developer portal platform:

### 1.12.1 Features

- **Software Catalog** - Track all services and resources
- **Templates** - Scaffold new projects
- **TechDocs** - Documentation as code
- **Plugins** - Extensible architecture

### 1.12.2 Catalog Entity

```
apiVersion: backstage.io/v1alpha1
kind: Component
metadata:
  name: my-service
  description: My microservice
  annotations:
    github.com/project-slug: myorg/my-service
    backstage.io/techdocs-ref: dir:.
spec:
  type: service
  lifecycle: production
  owner: team-platform
  system: my-system
  dependsOn:
  - resource:my-database
  providesApis:
  - my-api
```

### 1.12.3 Software Template

```
apiVersion: scaffolder.backstage.io/v1beta3
kind: Template
metadata:
  name: microservice-template
  title: Microservice Template
spec:
  owner: team-platform
  type: service
  parameters:
  - title: Service Details
    properties:
      name:
        type: string
```

```
    title: Service Name
  owner:
    type: string
    title: Owner Team
  steps:
    - id: fetch
      name: Fetch Template
      action: fetch:template
      input:
        url: ./skeleton
      values:
        name: ${{ parameters.name }}
    - id: publish
      name: Publish to GitHub
      action: publish:github
      input:
        repoUrl: github.com?owner=myorg&repo=${{ parameters.name }}
```

---

## 1.13 Platform Patterns

### 1.13.1 Golden Paths

Pre-defined, well-supported ways to accomplish common tasks:

1. **Service Creation** - Template with CI/CD, monitoring, logging
2. **Database Provisioning** - Self-service with guardrails
3. **Environment Management** - Dev, staging, production workflows

### 1.13.2 Self-Service Infrastructure

```
# Developer requests database
apiVersion: platform.example.com/v1
kind: Database
metadata:
  name: my-app-db
  namespace: my-team
spec:
  size: medium
  engine: postgres
```

### 1.13.3 Platform APIs

Expose platform capabilities through Kubernetes CRDs:

- Databases
- Message queues
- Storage buckets
- DNS entries

- Certificates
- 


## 1.14 Best Practices

1. **Start small** - Begin with highest-impact capabilities
  2. **Measure adoption** - Track developer satisfaction
  3. **Document everything** - Self-service requires good docs
  4. **Iterate based on feedback** - Platform is a product
  5. **Maintain golden paths** - Keep them up to date
- 

[← Back to CNPE Overview](#)

---

## 1.15 Sample Practice Questions



[Download PDF Version](#)

## 1.16 Practice Resources

- [Platform Engineering](#)
  - [Kubernetes Operators](#)
  - [Killercoda Scenarios](#)
- 

## 1.17 Platform Design (20%)

### 1.17.1 Question 1

Design a multi-tenant platform architecture for 50 development teams.

Show Solution

**Architecture considerations:**

#### 1. Isolation Strategy:

- Namespace per team with resource quotas
- Network policies for isolation
- RBAC per namespace

#### 2. Resource Management:

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: team-quota
  namespace: team-alpha
spec:
  hard:
    requests.cpu: "10"
    requests.memory: 20Gi
    limits.cpu: "20"
    limits.memory: 40Gi
    pods: "50"
```

#### 3. Shared Services:

- Centralized logging (Loki)
- Metrics (Prometheus/Thanos)
- Ingress controller
- Certificate management

#### 4. Self-Service:

- Namespace provisioning via GitOps
- Template-based deployments
- Developer portal (Backstage)

## 1.17.2 Question 2

How do you design APIs for a platform?

Show Solution

### Platform API Design Principles:

#### 1. Use Kubernetes-native APIs:

```
apiVersion: platform.example.com/v1
kind: Application
metadata:
  name: my-app
spec:
  image: myapp:v1
  replicas: 3
  ingress:
    enabled: true
    host: myapp.example.com
```

#### 2. Abstract complexity:

- Hide infrastructure details
- Provide sensible defaults
- Allow overrides when needed

#### 3. Versioning:

- Use API versioning (v1alpha1, v1beta1, v1)
- Maintain backward compatibility
- Deprecation policy

#### 4. Documentation:

- OpenAPI specs
  - Examples for common use cases
- 

## 1.18 Platform Implementation (30%)

### 1.18.1 Question 3

Create a Kubernetes Operator that manages a custom Application resource.

Show Solution

```
// Controller reconcile function
func (r *ApplicationReconciler) Reconcile(ctx
    context.Context, req ctrl.Request) (ctrl.Result, error) {
    var app platformv1.Application
    if err := r.Get(ctx, req.NamespacedName, &app); err != nil {
        return ctrl.Result{}, client.IgnoreNotFound(err)
    }
    // Reconcile logic here
    ...
    return ctrl.Result{}, nil
}
```

```

    }

    // Create Deployment
    deployment := &appsv1.Deployment{
        ObjectMeta: metav1.ObjectMeta{
            Name:      app.Name,
            Namespace: app.Namespace,
        },
        Spec: appsv1.DeploymentSpec{
            Replicas: &app.Spec.Replicas,
            Selector: &metav1.LabelSelector{
                MatchLabels: map[string]string{"app": app.Name},
            },
            Template: corev1.PodTemplateSpec{
                ObjectMeta: metav1.ObjectMeta{
                    Labels: map[string]string{"app": app.Name},
                },
                Spec: corev1.PodSpec{
                    Containers: []corev1.Container{
                        Name:  app.Name,
                        Image: app.Spec.Image,
                    },
                },
            },
        },
    }

    if err := ctrl.SetControllerReference(&app, deployment,
        r.Scheme); err != nil {
        return ctrl.Result{}, err
    }

    if err := r.Create(ctx, deployment); err != nil {
        if !errors.Exists(err) {
            return ctrl.Result{}, err
        }
    }

    return ctrl.Result{}, nil
}

```

## 1.18.2 Question 4

Implement GitOps for platform configuration.

Show Solution

### Repository Structure:

```

platform-config/
  └── clusters/
      └── production/
          └── flux-system/

```

```

    └── kustomization.yaml
    └── staging/
└── infrastructure/
    ├── controllers/
    ├── configs/
    └── kustomization.yaml
└── tenants/
    ├── team-alpha/
    │   ├── namespace.yaml
    │   ├── rbac.yaml
    │   └── quota.yaml
    └── team-beta/

```

### **Flux Kustomization:**

```

apiVersion: kustomize.toolkit.fluxcd.io/v1
kind: Kustomization
metadata:
  name: infrastructure
  namespace: flux-system
spec:
  interval: 10m
  sourceRef:
    kind: GitRepository
    name: platform-config
  path: ./infrastructure
  prune: true
  healthChecks:
  - apiVersion: apps/v1
    kind: Deployment
    name: ingress-nginx-controller
    namespace: ingress-nginx

```

### **1.18.3 Question 5**

Create a Crossplane Composition for provisioning databases.

Show Solution

```

apiVersion: apiextensions.crossplane.io/v1
kind: Composition
metadata:
  name: postgresql-aws
spec:
  compositeTypeRef:
    apiVersion: database.platform.io/v1
    kind: PostgreSQLInstance
  resources:
  - name: rds-instance
    base:
      apiVersion: rds.aws.crossplane.io/v1beta1
      kind: Instance
      spec:

```

```

forProvider:
  engine: postgres
  engineVersion: "14"
  instanceClass: db.t3.medium
  allocatedStorage: 20
  publiclyAccessible: false
patches:
- fromFieldPath: spec.storageGB
  toFieldPath: spec.forProvider.allocatedStorage
- fromFieldPath: spec.size
  toFieldPath: spec.forProvider.instanceClass
transforms:
- type: map
  map:
    small: db.t3.small
    medium: db.t3.medium
    large: db.t3.large
---
apiVersion: apiextensions.crossplane.io/v1
kind: CompositeResourceDefinition
metadata:
  name: postgresqlinstances.database.platform.io
spec:
  group: database.platform.io
  names:
    kind: PostgreSQLInstance
    plural: postgresqlinstances
  versions:
- name: v1
  served: true
  referenceable: true
  schema:
    openAPIV3Schema:
      type: object
      properties:
        spec:
          type: object
          properties:
            size:
              type: string
              enum: [small, medium, large]
            storageGB:
              type: integer

```

---

## 1.19 Platform Operations (25%)

### 1.19.1 Question 6

Implement a cluster upgrade strategy with zero downtime.

Show Solution

### Blue-Green Cluster Upgrade:

#### 1. Preparation:

```
# Create new cluster with updated version
eksctl create cluster -f new-cluster.yaml

# Install platform components
flux bootstrap github --context=new-cluster ...
```

#### 2. Migration:

```
# Sync workloads via GitOps
# Update DNS weights gradually

# Route 10% traffic to new cluster
aws route53 change-resource-record-sets \
--hosted-zone-id $ZONE_ID \
--change-batch file://weighted-10.json
```

#### 3. Validation:

```
# Monitor error rates
# Check application health
# Verify data consistency
```

#### 4. Cutover:

```
# Route 100% to new cluster
# Decommission old cluster
```

## 1.19.2 Question 7

Design a disaster recovery plan for a platform.

Show Solution

### DR Strategy:

#### 1. Backup:

```
# Velero backup schedule
apiVersion: velero.io/v1
kind: Schedule
metadata:
  name: daily-backup
spec:
  schedule: "0 2 * * *"
  template:
    includedNamespaces:
    - "*"
    excludedNamespaces:
```

```
- kube-system
storageLocation: aws-backup
ttl: 720h
```

## 2. Recovery Targets:

- RPO: 1 hour (data loss tolerance)
- RTO: 4 hours (recovery time)

## 3. DR Runbook:

- Restore cluster from backup
- Verify DNS failover
- Test application functionality
- Notify stakeholders

## 4. Regular Testing:

- Monthly DR drills
  - Documented procedures
  - Post-mortem reviews
- 

# 1.20 Security and Governance (15%)

## 1.20.1 Question 8

Implement policy-as-code for platform governance.

Show Solution

### Kyverno Policies:

```
# Require resource limits
apiVersion: kyverno.io/v1
kind: ClusterPolicy
metadata:
  name: require-limits
spec:
  validationFailureAction: Enforce
  rules:
    - name: check-limits
      match:
        any:
          - resources:
              kinds: [Pod]
      validate:
        message: "Resource limits required"
        pattern:
          spec:
            containers:
              - resources:
```

```

limits:
  memory: "?*"
  cpu: "?*"

---
# Require approved registries
apiVersion: kyverno.io/v1
kind: ClusterPolicy
metadata:
  name: allowed-registries
spec:
  validationFailureAction: Enforce
  rules:
    - name: check-registry
      match:
        any:
          - resources:
              kinds: [Pod]
      validate:
        message: "Images must be from approved registries"
        pattern:
          spec:
            containers:
              - image: "gcr.io/myorg/* | docker.io/myorg/*"

```

## 1.20.2 Question 9

Implement RBAC for platform teams.

Show Solution

```

# Platform Admin Role
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: platform-admin
rules:
  - apiGroups: ["*"]
    resources: ["*"]
    verbs: ["*"]

---
# Team Developer Role
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: team-developer
  namespace: team-alpha
rules:
  - apiGroups: [ "", "apps", "batch" ]
    resources: [ "pods", "deployments", "services", "configmaps",
                "secrets", "jobs" ]
    verbs: [ "get", "list", "watch", "create", "update", "delete" ]

```

```

- apiGroups: [""]
  resources: ["pods/log", "pods/exec"]
  verbs: ["get", "create"]

---
# Bind to team
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: team-alpha-developers
  namespace: team-alpha
subjects:
- kind: Group
  name: team-alpha-devs
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role
  name: team-developer
  apiGroup: rbac.authorization.k8s.io

```

---

## 1.21 Developer Experience (10%)

### 1.21.1 Question 10

Design a self-service namespace provisioning system.

Show Solution

**GitOps-based Provisioning:**

1. Request Template:

```

# teams/team-gamma/namespace-request.yaml
apiVersion: platform.io/v1
kind: TeamNamespace
metadata:
  name: team-gamma
spec:
  team: gamma
  owners:
  - user1@example.com
  - user2@example.com
  resourceQuota:
    cpu: "20"
    memory: 40Gi
  networkPolicy: restricted

```

2. Controller generates:

- Namespace
- ResourceQuota

- NetworkPolicy
- RBAC bindings
- Default LimitRange

### 3. Workflow:

- Developer submits PR
  - Automated validation
  - Approval from platform team
  - Merge triggers provisioning
- 

## 1.22 Exam Tips

1. **Know Kubernetes deeply** - Operators, controllers, CRDs
  2. **Practice GitOps** - Flux, Argo CD configurations
  3. **Understand IaC** - Crossplane, Terraform
  4. **Know policy tools** - Kyverno, OPA/Gatekeeper
  5. **Practice troubleshooting** - Logs, events, debugging
- 

[← Back to CNPE Overview](#)

---